


2020-11-05

└─What is smallpt anyway?

1. **S**
2. 99 LoC C++: **small** path tracer.
3. Ported to many languages, including Haskell!
4. Haskell port was by Vo Minh Thu. Thanks a ton!
5. Start from `noteed`'s original source; SHA the output image from the Haskell source for baseline.
6. Perfect for an optimization case study.
7. Plan: Quick walk through Haskell code, end up at C++ (`clang++`) performance.

What is smallpt anyway?

What is smallpt anyway?




- 99 LoC C++ **small** path tracer.
- Ported to many languages, including Haskell! (Thanks to Vo Minh Thu/noteed).
- Start from noteed's original source; SHA the output image from the Haskell source for baseline.

1. **S**
2. 99 LoC C++: **small** path tracer.
3. Ported to many languages, including Haskell!
4. Haskell port was by Vo Minh Thu. Thanks a ton!
5. Start from noteed's original source; SHA the output image from the Haskell source for baseline.
6. Perfect for an optimization case study.
7. Plan: Quick walk through Haskell code, end up at C++ (clang++) performance.

2020-11-05

└ A TL;DR of a path tracer



1. **S**
2. Brief spiritually correct description of how a path tracer works
3. Main problem: what light ray hits the eye?
4. Idea: trace backwards; start from the eye, hypothesize light came from a direction
5. Follow the direction, and see if light did indeed come from this direction
6. Hypothesize light came from direction r_1 . Follow and see what happens
7. Similarly for r_2 , r_3

2020-11-05

└ A TL;DR of a path tracer



1. Let's say our ray hits a light source
2. Then we know that the ray came from the light source
3. Set color to color of light source


2020-11-05

└ A TL;DR of a path tracer

VOID \longleftrightarrow D
black(0,0,0)

1. Let's say our ray hits nothing
2. Then we know that nothing could have produced this ray.
3. Set color to zero

└ A TL;DR of a path tracer




1. Let's say our ray hit a metallic object. This is neither a light source, nor nothing
2. We want to find light rays, which on striking the metal, produce our black right ray
3. Use reflection: angle of incidence equals angle of reflection
4. Perform math, find light ray that lead to black right ray
5. candidate blue ray is shown
6. recurse

Optimizing smallpt


2020-11-05

└ A TL;DR of a path tracer



1. Let's say our ray hit a glass object. This is different from all the previous cases
2. Here, refraction comes into play
3. Perform math, find light ray that lead to black right ray
4. candidate blue ray is shown
5. recurse

└ A TL;DR of a path tracer



1. Consider a difficult scene like this one, where light can only enter from the top
2. Light may need to bounce many times before it enters the eye
3. How many bounces do we consider?
4. Make longer bounces more unlikely
5. Setup a russian roulette system, where the longer a ray has bounced, the more likely it is to die (stop recursing)
6. increase number of bullets in the gun as number of bounces increase

What is smallpt anyway?

What is smallpt anyway?

```
struct Vec {
    double x, y, z; // position, also color (r,g,b)
    ... methods ...
};

struct Ray { Vec o, d; Ray(Vec o_, Vec d_) : o(o_), d(d_) {} };
enum Refl_t { DIFF, SPEC, REFR }; // material types, used in radiance()

struct Sphere {
    double radi; // radius
    Vec p, n, e; // position, emission, color
    Refl_t refl; // reflection type (DIFFuse, SPECular, REFRactive)
    ... methods ...
    double intersect(const Ray &r) const { // returns distance, 0 if miss
    };
    double sphererad[] = { //same: radius, position, emission, color, material
    Sphere(int i, Vec( i+0i,0i,0i+0i), Vec(1,Vec(1,2i,2i),DIFF), //left
    ... initialization ...
    );
};
```

1. S
2. Has geometric primitives: vectors, spheres, materials
3. Entirely numeric-based, no real “data structures” to speak of

└─What is smallpt anyway?



1. **S**
2. Most of the compute cost is spent in the function that traces rays.
3. is called radiance

What is smallest anyway?

```
var radiance = color; Ray ray; int depth; unsigned short *Xi[C];
```

radiance
radiance

radiance

radiance

radiance radiance
radiance radiance

}

1. **S**
2. radiance is the function that performs this path tracing
3. Recursively calls itself a bunch of times

What is smallpt anyway?



1. **S**
2. Recursion is guarded by a lot of control flow

—What is smallpt anyway?

```
What is smallest anyway?

var realisation = function (logLik, logPrior, logPosterior, nData) {

  // var = var(x), so the obj is named(), will be directly obj["var"] = f(x, y) obj["x"]
  if (
    ) {
      } else {

    } else {

      realisation
    }
    realisation

  if (is.null(x), newParam = 1) obj = obj["x"]
    realisation

  realisation
  realisation
  realisation
}
```

1. **S**
2. The control flow and computation is very numeric in nature

└─What is smallpt anyway?

[illegible]

1. **S**
2. uses the function `erand48` for randomness

—What is smallpt anyway?

What is smallpt anyway?

```

Vec randomVec() { Vec r; r.x = rand(); r.y = rand(); return r; }

double x; // Distance to information
double y; // Distance to destination
double z; // Distance to obstacle
double w; // Distance to goal
if (fabs(r.x - x) > 1) r.x = x; // If x is too far, return back
else { r.x = x + (rand() - RAND_MAX) * 0.1; } // else just step
if (fabs(r.y - y) > 1) r.y = y; // If y is too far, return back
else { r.y = y + (rand() - RAND_MAX) * 0.1; } // else just step
if (fabs(r.z - z) > 1) r.z = z; // If z is too far, return back
else { r.z = z + (rand() - RAND_MAX) * 0.1; } // else just step
if (fabs(r.w - w) > 1) r.w = w; // If w is too far, return back
else { r.w = w + (rand() - RAND_MAX) * 0.1; } // else just step
Vec v = (Vec)(r.x - x, r.y - y, r.z - z, r.w - w); // Vector from current
// to goal
double d = v.x * v.x + v.y * v.y + v.z * v.z + v.w * v.w; // Distance
// squared
if (d < 0.1) return r; // If distance is small, return goal
if (fabs(d) > 1000) return r; // If distance is large, return goal
Vec r2 = randomVec(); // Random vector
Vec v2 = (Vec)(r2.x - x, r2.y - y, r2.z - z, r2.w - w); // Vector from
// current to random vector
double d2 = v2.x * v2.x + v2.y * v2.y + v2.z * v2.z + v2.w * v2.w; // Distance
// squared
if (d2 < d) return r2; // If distance is smaller, return random vector
if (fabs(d2) > 1000) return r2; // If distance is large, return random vector
return r; // Return current vector
}

Vec solve() { Vec start = Vec(0, 0, 0, 0); // Start vector
Vec goal = Vec(10, 10, 10, 10); // Goal vector
Vec r = randomVec(); // Random vector
Vec v = (Vec)(r.x - start.x, r.y - start.y, r.z - start.z, r.w - start.w); // Vector from
// start to random vector
double d = v.x * v.x + v.y * v.y + v.z * v.z + v.w * v.w; // Distance
// squared
if (d < 0.1) return r; // If distance is small, return goal
if (fabs(d) > 1000) return r; // If distance is large, return goal
Vec r2 = randomVec(); // Random vector
Vec v2 = (Vec)(r2.x - start.x, r2.y - start.y, r2.z - start.z, r2.w - start.w); // Vector from
// start to random vector
double d2 = v2.x * v2.x + v2.y * v2.y + v2.z * v2.z + v2.w * v2.w; // Distance
// squared
if (d2 < d) return r2; // If distance is smaller, return random vector
if (fabs(d2) > 1000) return r2; // If distance is large, return goal
return r; // Return current vector
}

```

1. **S**
2. The full code continues to be more of the same

Initial Haskell Code: radiance (1x)

1. **S**
2. The same computation, this time in haskell

[illegible]

Initial Haskell Code: Data structures (1x)

Initial Haskell Code: Data structures (1x)

```

data Vec = Vec {-# UNPACK #-} !Double {-# UNPACK #-} !Double {-# UNPACK #-} !Double

norm :: Vec -> Vec -> Vec
[-1] :: Vec -> Double -> Vec
radial f =
  let
    norm :: Vec -> Double
    norm :: Vec -> Vec
    norm v = v * norm (norm v)
    dot :: Vec -> Vec -> Double
    cross :: Vec -> Double
  in
    norm

data Ray = Ray Vec Vec -- origin, direction
data Refl = DIFF | SPEC | REFR -- material type, used in radiance
-- | radius, position, emission, color, refraction
data Sphere = Sphere Double Vec Vec Vec Refl

```

1. S
2. We implement the same geometric data structures in Haskell
3. Code here has an inconsistency
4. while Vec has unpack, Ray, Sphere not having unpack

Initial Haskell code: Sphere intersection

Initial Haskell code: Sphere intersection

```

intersect :: Ray -> Sphere -> Maybe Double
intersect (Ray a d) (Sphere r p _a _d _refl) =
  if det < 0 then Nothing else f (b-side) (b+side)
  where op = p - a -- Vector
        b = dot op d
        det = b*b - dot op op - r*r -- Numeric
        side = sqrt det
        f a = if a < op then Just a else if a > op then Just a else Nothing
intersect :: Ray -> (Maybe Double, Sphere)
intersect ray = (h, s)
  where (h,s) = foldl' f (Nothing,undefined) spheres -- Spheres iterated over
        f (h',sp) a' = case (h',intersect ray a') of
          (Nothing,Just a) -> (Just a,a')
          (Just p,Just a) | a < p -> (Just a,a')
          _ -> (h',sp)

```

1. S
2. Responsible for figuring out what the ray hits.
3. We iterate over the list of spheres.
4. Once again, numeric heavy.
5. Use a Maybe to indicate whether we've found an answer or not.

Initial Haskell Code: radiance (1x)

[illegible]

1. **S**
2. Branch heavy
3. Recursive
4. Uses an RNG

Initial Haskell Code: Entry point (1×)

```
Initial Haskell Code: Entry point (1×)

mainPrt = let -> let -> let -> IO ()
mainPrt = h sample -> do
  ...
  d <- VM.replicate (x + y) 0
  allocateArray 0 (x+1) -> Create mutable memory
  flip mapM [0..h-1] $ \y -> do -- loop
    writeM xi y
    for_ [0..w-1] \x -> do -- loop
      let s = (h-y-1) * w + x
      for_ [0..1] \ay -> do -- loop
        r <- readMRef 0 -- Create mutable memory
        for_ [0..sample-1] \a -> do -- loop, loopa
          r1 <- (do) (do) readMRef xi
          ...
          rad <- radiance (Ray (Ray-d+100) (fromIntegral s)) 0 xi -- branch
          ...
          modifyMRef r (- rad -> recip (fromIntegral sample)) -- write
          ci <- VM.unsafeRead c i
          let r1 r1 <- readMRef r
          VM.unsafeWrite c i $
            ci + (do (clamp r1) (clamp r1) (clamp r1) -> 0.25 -- write
          ...
  ...
```

1. S
2. Use mutability to store the pixels of the image in c
3. Loops over all pixels in an image and shoots rays
4. Shoots sample number of rays per pixel and adds up the results
5. Finally, writes resulting color out by mutating the array c

Initial Haskell Code: RNG (1×)

```
Initial Haskell Code: RNG (1×)  
  
foreignImport ccall unsafe "erand48"  
erand48 :: Ptr CDouble -> IO Double
```

1. **S**
2. As mentioned previously, we use the RNG to decide randomly in which direction to send rays
3. erand48 is imported as a foreign ccall for parity with the C code

└ Restrict export list to main ($1\times \mapsto 1.13\times$)

Restrict export list to main ($1\times \mapsto 1.13\times$)

```
-module Main where  
module Main (main) where
```

1. **S**
2. The very first thing to do is to let the compiler actually optimize.
3. If a function is public, then compiler doesn't know all call sites
4. Export only the one function that's called from the outside: main
5. Allows compiler to know that other functions in module are not called from outside
6. compiler has "perfect knowledge" about these functions now

└ Restrict export list to main ($1\times \mapsto 1.13\times$)

Restrict export list to main ($1\times \mapsto 1.13\times$)

```
-module Main where  
module Main (main) where  
  ▶ Exported functions could be used by something unknown.  
  ▶ Original versions must be available.
```

1. **S**
2. The very first thing to do is to let the compiler actually optimize.
3. If a function is public, then compiler doesn't know all call sites
4. Export only the one function that's called from the outside: main
5. Allows compiler to know that other functions in module are not called from outside
6. compiler has "perfect knowledge" about these functions now

└─ Mark entries of Ray and Sphere as UNPACK and Strict
(1.13x \rightarrow 1.07x)

```
data Vec = Vec {#d UNPACK d-} | Double
  (#d UNPACK d-} | Double
  (#d UNPACK d-} | Double)

data Ray = Ray Vec Vec --> origin, direction
data Ray = Ray Vec Vec --> origin, direction

data Sph = SPH {#d UNPACK d-} | Double
  (#d UNPACK d-} | Double
  (#d UNPACK d-} | Double)

-- radius, position, emission, color, and function
data Sphere = Sphere Double Vec Vec Vec Sph
  (#d UNPACK d-} | Double
  (#d UNPACK d-} | Vec
  (#d UNPACK d-} | Vec
  (#d UNPACK d-} | Vec Sph)

strict Vec {#d UNPACK d-}
strict Ray {#d UNPACK d-}
strict Sphere {#d UNPACK d-}
strict RaySph {#d UNPACK d-}
strict RaySph {#d UNPACK d-}
strict RaySph {#d UNPACK d-}
strict RaySph {#d UNPACK d-}
```

1. D
2. Strictness in the arguments means that they're evaluated when instantiated, not when demanded.
3. Where as Unpacking removes indirection from doing a memory lookup for components.
4. Means we have to copy everything into the data structure that it is unpacked into.
5. We don't unpack ray (Lots of calculations on its components, want those to fuse)
6. Unpack Sphere - its static from compile time
7. Don't unpack Ray, because each Vec undergoes a lot of computation.

Use a pattern synonym to unpack Refl in Sphere
($1.07 \times \mapsto 1.07 \times$)

Use a pattern synonym to unpack Refl in Sphere ($1.07 \times \mapsto 1.07 \times$)

```

+ (4 LAMBDA2 PatternSynonym R)
data Refl = DIFF | SPEC | REFR -- material types, used in radiance
newtype Refl' = Refl' Refl -- material types, used in radiance
pattern DIFF, SPEC :: Refl
pattern DIFF = Refl 0
pattern SPEC = Refl 1
pattern REFR = Refl 2
+ (4 COMPLETE DIFF, SPEC, REFR R)

-- radius, position, emission, color, refraction
data Sphere = Sphere {4 DIFF2 R} Double
    {4 DIFF2 R} RVec
    {4 DIFF2 R} RVec
    {4 DIFF2 R} RVec RVec1
+
+   {4 DIFF2 R} RVec {4 DIFF2 R} RVec1

```

1. D
2. Was unable to unpack Refl
3. UnboxedSums are recent
4. UnboxedSums are very unpleasant
5. We're using an older trick to fake the unboxing here instead.
6. In this case it isn't much of a win, but it illustrates the technique.

Change from maximum on a list to max ($1.07\times \mapsto 1.08\times$)

```
Change from maximum on a list to max (1.07x → 1.08x)

--max (Vec a b c) = maximum [a,b,c]
--max (Vec a b c) = max a (max b c)

let x = a "add" (d "mul" t)
a = max b a "add" p
all = if a "add" d < 0 then a else a "mul" (-d)
--
p = max c
depth = depth + 1
continue f = case refi of
  RAY -> do
...
if depth > 0
then do
  w <- randomM 0 1
  let tpe = max a
+

```

1. D
2. Prebuild comparison
3. Don't go via list
4. GHC does not evaluate at compile time, only has RULES
5. Doesn't really help much in this case

- Convert `erand48` to pure Haskell ($1.08\times \mapsto 1.10\times$)

[illegible]

1. **S**
2. The entire premise of this talk is that Haskell can be as fast as C.
3. We're opening the black box of what `erand48` does to GHC
4. Further any impedance mismatch, such as FFI almost universally has to have, carries some bookkeeping overhead.
5. If our Haskell code was as fast as the C code moving the code into Haskell would be a win, if it was slightly slower it could still be a win.
6. Often considering your Haskell code's performance is a better option and easier than reimplementing something in C.
7. As is the way with optimizations, this is not universally true.

- Remove mutability: Erand48 Monad ($1.10\times \mapsto 1.15\times$)

1. **D**
2. All these mutability locations throw in extra RTS code, extra sequencing that blocks the compiler's optimization, and dependency chains.
3. Sometimes we need mutability for performance.
4. SSA is normal to compilers though.
5. We almost start at SSA as a functional language.
6. don't break it when you don't have a good reason.

L

[illegible]

1. D

- Set **everything** in smallpt to be strict ($1.10\times \mapsto 1.15\times$)

Set everything in `smallpt` to be strict (1.10x \leftrightarrow 1.15x)

[illegible]

Don't senselessly bang everything in sight.

1. **D**
2. This is not a recommendation, this is a warning.
3. We get a speedup here but it can also regress performance. Some of these bangs are regressions that are hidden.

Why strictness may be bad

Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
```

1. **S**
2. Consider the function `foo` and `fooOpt`. These are equivalent
3. The fact that `x` is not used allows us to eliminate computing `x`
4. Consider the next version
5. Illegal, we need to have `x`, because it doesn't produce `ERR`
6. we can't equationally reason about the program anymore.
7. Makes it harder for `GHC`. `GHC` is conservative about bangs
8. Inhibits compiler from optimizing

2020-11-05

Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
Prelude> let foo = let x = error "ERR" in \y -> y
Prelude> foo 12
12
```

1. **S**
2. Consider the function `foo` and `fooOpt`. These are equivalent
3. The fact that `x` is not used allows us to eliminate computing `x`
4. Consider the next version
5. Illegal, we need to have `x`, because it doesn't produce `ERR`
6. we can't equationally reason about the program anymore.
7. Makes it harder for `GHC`. `GHC` is conservative about bangs
8. Inhibits compiler from optimizing

2020-11-05

Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
Prelude> let foo = let x = error "ERR" in \y -> y
Prelude> foo 12
12
let fooOpt = \y -> y
```

1. **S**
2. Consider the function `foo` and `fooOpt`. These are equivalent
3. The fact that `x` is not used allows us to eliminate computing `x`
4. Consider the next version
5. Illegal, we need to have `x`, because it doesn't produce `ERR`
6. we can't equationally reason about the program anymore.
7. Makes it harder for `GHC`. `GHC` is conservative about bangs
8. Inhibits compiler from optimizing

Why strictness may be bad

Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
Prelude> let foo = let x = error "ERR" in \y -> y
Prelude> foo 12
12
let fooOpt = \y -> y
let foo' = let !x = error "ERR" in \y -> y
```

1. **S**
2. Consider the function `foo` and `fooOpt`. These are equivalent
3. The fact that `x` is not used allows us to eliminate computing `x`
4. Consider the next version
5. Illegal, we need to have `x`, because it doesn't produce `ERR`
6. we can't equationally reason about the program anymore.
7. Makes it harder for `GHC`. `GHC` is conservative about bangs
8. Inhibits compiler from optimizing

Why strictness may be bad

Why strictness may be bad

```

let foo = let x = error "ERR" in \y -> y
Prelude> let foo = let x = error "ERR" in \y -> y
Prelude> foo 12
12
let fooOpt = \y -> y
let foo' = let !x = error "ERR" in \y -> y
Prelude> let foo' = let !x = error "ERR" in \y -> y
Prelude> foo' 12
*** Exception: ERR
CallStack (from HasCallStack):
  error, called at <interactive>:0:21 in interactive:Ghci2

```

1. S
2. Consider the function foo and fooOpt. These are equivalent
3. The fact that x is not used allows us to eliminate computing x
4. Consider the next version
5. Illegal, we need to have x, because it doesn't produce ERR
6. we can't equationally reason about the program anymore.
7. Makes it harder for GHC. GHC is conservative about bangs
8. Inhibits compiler from optimizing

Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
Prelude> let foo = let x = error "ERR" in \y -> y
Prelude> foo 12
12
let fooOpt = \y -> y
let foo' = let !x = error "ERR" in \y -> y
Prelude> let foo' = let !x = error "ERR" in \y -> y
Prelude> foo' 12
*** Exception: ERR
CallStack (from HasCallStack):
  error, called at <interactive>:8:21 in interactive:Ghci2
let fooOpt = \y -> y --> INTERPRETING: forcing foo'~fooOpt should give "ERR"
```

1. S
2. Consider the function foo and fooOpt. These are equivalent
3. The fact that x is not used allows us to eliminate computing x
4. Consider the next version
5. Illegal, we need to have x, because it doesn't produce ERR
6. we can't equationally reason about the program anymore.
7. Makes it harder for GHC. GHC is conservative about bangs
8. Inhibits compiler from optimizing

- Reduce to only useful strictnesses in smallpt
(1.10x \mapsto 1.15x)

1. **D**
2. Thus the compiler can no longer move the computation around or simplify it.
3. Force useless work.
4. A little thinking about how the variables are used or looking at core allows us to select which ones we bang selectively.

└ Strategic application of strictness in entire project
 (1.15x \mapsto 1.23x)

Strategic application of strictness in entire project (1.15x \mapsto 1.23x)

```

...
+ if desc0 then RayTracing strict 0 (bounces) (bounces)
+
+ where up = p - o
+     ngs = (n · u)
+
+     b = desc0 up · d
+     den = 1/b^2 // den up = up
+     nlen = ngs den
+     f n n = if ngs then desc0 n nlen if ngs then desc0 n nlen RayTracing
+
+ if desc0
+   then RayTracing
+   else
+     let ngs = (n · d)
+     let den = 1/ngs den
+     let f n = bounces
+     let f n = bounces
+     let f n = bounces
+
+   in if ngs then desc0 n nlen if ngs then desc0 n nlen RayTracing
+
+
  ...

```

1. D
2. Sometimes (point out 'intersect') we have to rearrange the code though when we use bangs.
3. Bangs tell the compiler to make more efficient code, but take away the compiler's options in how to do so.
4. Only take away the compiler's liberties when it's using them poorly.
5. Becomes intuitive.

- Remove Maybe from intersect(s) ($1.23\times \mapsto 1.40\times$)

[illegible]

1. **S**
2. This is a far more performance critical version of what we saw with 'maximum' vs. 'max'.
3. innermost functions are of critical importance. remove Maybe which significantly reduces the boxing
4. Since a Ray that fails to intersect something can be said to intersect at infinity, Double already actually covers the structure at play
5. This also reduces allocation.

Hand unroll the fold in intersects ($1.40\times \mapsto 1.43\times$)

Hand unroll the fold in intersects ($1.40x \mapsto 1.43x$)

```

intersectum := Ray → {Double, Sphere}
intersectum ray = {x,y,z}
  where (x,y,z) = find (λ (t,d) → t/2.0 <= d.2) spheres
intersectum ray =
  if ... (if (intersect ray sphLeft, sphLeft) sphRight) ...
  where
    if (x',y,z) = let h ← intersect ray sph in if x' < h then (x,y,z) else (x',y,z)

spheres := {Sphere}
spheres = let x = Sphere x ← a ← aVec → (a ← mVec) → y ← Vec in
  [ let x = Sphere (Vec (Vec (Vec 0.0 0.0 0.4) (Vec (Vec 0.78 0.28 0.28) 0.02)) → Left
    , let x = Vec (Vec (Vec 0.0 0.0 0.4) (Vec (Vec 0.28 0.28 0.78) 0.02)) → Right
  ]

sphLeft, sphRight, ... :: Sphere
sphLeft = Sphere (Vec (Vec (Vec 0.0 0.0 0.4) aVec) (Vec 0.78 0.28 0.28) 0.02)
sphRight = Sphere (Vec (Vec (Vec 0.0 0.0 0.4) aVec) (Vec 0.28 0.28 0.78) 0.02)
...

```

1. **D**
2. 'intersects' is very hot
3. Loop unrolling
4. Many compilers do this for us, and there are special versions of it like Duff's Device.
5. Sadly GHC doesn't
6. Can do variants by hand.
7. RULE could handle each one specifically (only exactly that one)?

Custom datatype for intersects parameter passing ($1.43 \times \mapsto 1.46 \times$)

```

Custom datatype for intersects parameter passing (1.43x  $\mapsto$  1.46x)

Std: Tuple with possibly-uninitialized Double and Sphere
Req: Reference to a guaranteed-to-be-evaluated Double and Sphere
intersects :: Ray -> (Double, Sphere)
- data T = T (Double) (Sphere)
+
+ intersects :: Ray -> T
+ intersects ray =
+   if (...) T (intersect ray sphLeft, sphLeft) sphRight) ... sphLeft
+   where
+     - if (s', sp) s' =
+     -   let ts = intersect ray s' in if x < ts then (ts, s') else (s', sp)
+     - if (t', sp) ts =
+     -   let ts = intersect ray s' in if x < ts then T ts s' else T s' sp
+
+ radiuses :: Ray -> Int -> Random8 Vec
+ radiuses rpfRay n d1 depth = case intersects ray of
+   (ts,_) | t == 1/0.0 -> return 0
+   (ts, (Sphere _ p s rad)) -> do
+     [t',_] | t == 1/0.0 -> return 0
+     [t', (Sphere _ p s rad)] -> do
+       let ts = s - d * s
+       ts = norm 0 ts - p
+       ts = if dot s d < 0 then s else negate s

```

1. D
2. We can optimize data passing.
3. Want: Data strict, but not unpacked.
4. Compiler knows its evaluated but no copying
5. A normal tuple lacks strictness information.
6. An unboxed tuple forces copying
7. Strict Tuple.
8. This exists in libraries of course, but we wanted to illustrate it.

└─ Optimize file writing: $(1.46\times \mapsto 1.46\times)$

Optimize file writing: $(1.46\times \mapsto 1.46\times)$

```
build-dependencies:
  base >= 4.12 && < 4.15
+ , bytestring >= 0.11
+
-totat :: Double -> Int
+totat :: Double -> Int
-totat x = floor $ clamp x ** recip 2.2 * 255 + 0.5
+totat :: Double -> BB.Builder -- O(1) concatenation
-totat x = BB.intDec (floor (clamp x ** recip 2.2 * 255 + 0.5)) <> BB.char8 ' '
+...
withFile "image.ppm" WriteMode $ \hdl -> do
-  hPutStrLn hdl "P3\n# Grad3\n" < b 255 :: Int)
-  for_ img \fVec x g b -> do
-    hPutStrLn hdl "255\n" (totat x) (totat g) (totat b)
-    BB.BuilderBuilder hdl $
+  BB.BuilderBuilder hdl $
+    BB.intDec x <> BB.char8 ' ' <> BB.intDec g <> BB.char8 'a' <>
+    BB.intDec 255 <> BB.char8 'a' <>
+    (concat $ map \fVec x g b -> totat x <> totat g <> totat b) img)
```

1. D
2. Strings are inefficient
3. 'bytestring' has some efficient writing code, so we just convert to that for a modest gain.

└ Use LLVM backend ($1.46\times \mapsto 2.04\times$)

```
Use LLVM backend (1.46x ==> 2.04x)


package smallpt-opt
+ ghc-options: -fllvm
```

1. **S**
2. Finally, this particular code is quite numeric heavy.
3. There are optimizations for numeric heavy code we're missing in GHC.
4. LLVM has an extensive library of laws to optimize low level numeric ops.
5. LLVM is too low-level to understand haskell as haskell.
6. LLVM makes decisions with the tacit assumption that the assembly came from a C-like language, which is often to the detriment of a Haskell-like language.
7. In this case, as the code is "fortran-like", LLVM wins.

└ The view from the mountaintop ($1.46\times \mapsto 2.04\times$)

1. D

The view from the mountaintop ($1.46\times \mapsto 2.04\times$)



└─ Avoid CPU iieee754 slow paths ($2.04\times \mapsto 2.12\times$)

```

Avoid CPU iieee754 slow paths (2.04x  $\mapsto$  2.12x)
C source:
inline bool intersect(...) {
    ...
    iieee754;
}
return t>inf;
}
Improvement:
intersect :: Ray -> Sphere -> Double
intersect (Ray a d) (Sphere s p _s _refl) =
    if d>0
    + then 1/0.0
    + else 1e20
else
    ...
-   in if s>eps then a else if s>eps then a else 1/0.0
+   in if s>eps then a else if s>eps then a else 1e20
...
radiance :: Ray -> Int -> Random00 Vec
radiance ray(Ray a d) depth = case intersect ray of
    + [ (t, s) ] | t <= 1e20.0 -> return 0
    + [ (t, s0) ] -> return 0
    ...

```

1. D
2. We used +Inf to match the Maybeness
3. C++ code set 1e20 s the horizon
4. Mechanical sympathy is important.
5. Know how the CPU (abstractly) executes - slow path / fast path.

└ Fix differences with C++ version ($2.12\times \mapsto 2.32\times$)

```
Fix differences with C++ version (2.12x  $\mapsto$  2.32x)
10x10 image size, at commit 16d5841: the LLVM backend,
1020d2479baee7738a2080079c1400a103020e clang++ -ppm
92d1d424e0898677e05a08400457075c50a12e4d g++ -ppm
4862a76d11670f705c23077e4a3d0e004a3a111089 gbc -ppm
```

1. S
2. Since the sha1 of the output didn't match the C++ version we started investigating.
3. clang++, g++ actually produce different sha1s
4. unincremented depth was being used in one branch, causing us to do more work
5. now confident to say we're doing the same computation as C++

└ Fix differences with C++ version ($2.12\times \mapsto 2.32\times$)

```
Fix differences with C++ version (2.12x  $\mapsto$  2.32x)

10x10 image size, at compile. SHA1s: Use LLVM backend,
1020f2d79baee773ba080079c440ba03020c clang++ .ppm
9261d424e080677e04a0400457075c5da32e62 g++ .ppm
6b62a7a11cf28705c2307e4a0b0040a111089 ghc .ppm

Subtle mutation:
if (--depth < 0)
...
return ... + (depth < 0 ? ... : ...) // depth is after depth++
```

1. S
2. Since the sha1 of the output didn't match the C++ version we started investigating.
3. clang++, g++ actually produce different sha1s
4. unincremented depth was being used in one branch, causing us to do more work
5. now confident to say we're doing the same computation as C++

└ Fix differences with C++ version ($2.12\times \mapsto 2.32\times$)

Fix differences with C++ version ($2.12\times \mapsto 2.32\times$)

```

$D1D0 image size, at commit: 56d8d61 Use LLVM backend,
1020f2d79aee7738a080076c4408e0303c clang++ -pp
9261d424e080677e04a080004707550a32e02 g++ -pp
6b62a7a1d1f28703c3207e4a080000a0a11080 ghc -pp

Subtle mutation:
if (--depth>0)
...
return ... + (depth>0 ? ... : ...) // depth is after depth++
Fix:
let depth' = depth + 1
in
...
-
+           if depth>0
+           if depth'>0 -- depth' = depth + 1
...

```

1. S
2. Since the sha1 of the output didn't match the C++ version we started investigating.
3. clang++, g++ actually produce different sha1s
4. unincremented depth was being used in one branch, causing us to do more work
5. now confident to say we're doing the same computation as C++

└ Fix differences with C++ version ($2.12\times \mapsto 2.32\times$)

Fix differences with C++ version ($2.12\times \mapsto 2.32\times$)

```

10x10 image size, at commit 16d661: Use LLVM backend.
1020d2d7b4ee67138d288075c1e10be1232dfe clang++.ppm
9261d42e6898677e03a08480647075c5d932e4d g++.ppm
68b2a7611cf187f55c12307e5a3d88684b111089 ghc.ppm

Subtle mutation:
if (--depth < 0)
...
return ... + (depth < 0 ? ... : ...) // depth is after depth++
Fix:
let depth' = depth + 1
...
is
...
-           if depth < 0
+           if depth' < 0 -- depth' = depth + 1
...

After:
10x10 image size, at commit 16d661: Fix differences with C++ version
1020d2d7b4ee67138d288075c1e10be1232dfe clang++.ppm
9261d42e6898677e03a08480647075c5d932e4d g++.ppm
1020d2d7b4ee67138d288075c1e10be1232dfe ghc.ppm

```

1. S
2. Since the sha1 of the output didn't match the C++ version we started investigating.
3. clang++, g++ actually produce different sha1s
4. unincremented depth was being used in one branch, causing us to do more work
5. now confident to say we're doing the same computation as C++

Takeaways

Takeaways

- ▶ The unrolling in 'intersect' is ugly
- ▶ (We feel) the maintainability of this code hasn't been significantly harmed.
- ▶ We're faster than clang++ and g++
- ▶ Haven't exhausted the optimization opportunities.
- ▶ SSE could learn to do several of these optimizations for us.
- ▶ Others are just good Haskell style.
- ▶ Clean Haskell is often performant Haskell.
- ▶ Repository stepping through each optimization is available at github.com/keel1n/smallpt-opt
- ▶ Slides at github.com/keel1n/wides-haskell-exchange-2020-venuept

1. D