


2020-11-05

└─What is smallpt anyway?

1. **S**
2. 99 LoC C++: **small** path tracer.
3. Ported to many languages, including Haskell!
4. Haskell port was by Vo Minh Thu. Thanks a ton!
5. Start from `noteed`'s original source; SHA the output image from the Haskell source for baseline.
6. Perfect for an optimization case study.
7. Plan: Quick walk through Haskell code, end up at C++ (`clang++`) performance.

2020-11-05

What is smallpt anyway?




- 99 LoC C++ **small** path tracer.
- Ported to many languages, including Haskell! (Thanks to Vo Minh Thu/noteed).
- Start from noteed's original source; SHA the output image from the Haskell source for baseline.

1. **S**
2. 99 LoC C++: **small** path tracer.
3. Ported to many languages, including Haskell!
4. Haskell port was by Vo Minh Thu. Thanks a ton!
5. Start from noteed's original source; SHA the output image from the Haskell source for baseline.
6. Perfect for an optimization case study.
7. Plan: Quick walk through Haskell code, end up at C++ (clang++) performance.

2020-11-05

└ A TL;DR of a path tracer



1. **S**
2. Brief spiritually correct description of how a path tracer works
3. Main problem: what light ray hits the eye?
4. Idea: trace backwards; start from the eye, hypothesize light came from a direction
5. Follow the direction, and see if light did indeed come from this direction
6. Hypothesize light came from direction  $r_1$ . Follow and see what happens
7. Similarly for  $r_2$ ,  $r_3$

2020-11-05

└ A TL;DR of a path tracer



1. Let's say our ray hits a light source
2. Then we know that the ray came from the light source
3. Set color to color of light source

## Optimizing smallpt


2020-11-05

└ A TL;DR of a path tracer

VOID  $\longleftrightarrow$  D  
black(0,0,0)

1. Let's say our ray hits nothing
2. Then we know that nothing could have produced this ray.
3. Set color to zero

└ A TL;DR of a path tracer




1. Let's say our ray hit a metallic object. This is neither a light source, nor nothing
2. We want to find light rays, which on striking the metal, produce our black right ray
3. Use reflection: angle of incidence equals angle of reflection
4. Perform math, find light ray that lead to black right ray
5. candidate blue ray is shown
6. recurse

## Optimizing smallpt

2020-11-05

└ A TL;DR of a path tracer



1. Let's say our ray hit a glass object. This is different from all the previous cases
2. Here, refraction comes into play
3. Perform math, find light ray that lead to black right ray
4. candidate blue ray is shown
5. recurse

└ A TL;DR of a path tracer



► When to end recursion?  
► Choose to end recursion randomly after some large number of rounds  
► "Russian Roulette"

1. Consider a difficult scene like this one, where light can only enter from the top
2. Light may need to bounce many times before it enters the eye
3. How many bounces do we consider?
4. Make longer bounces more unlikely
5. Setup a russian roulette system, where the longer a ray has bounced, the more likely it is to die (stop recursing)
6. increase number of bullets in the gun as number of bounces increase



What is smallpt anyway?

What is smallpt anyway?

```
struct Vec {
    double x, y, z; // position, also color (r,g,b)
    ... methods ...
};

struct Ray { Vec o, d; Ray(Vec o_, Vec d_) : o(o_), d(d_) {} };
enum Refl_t { DIFF, SPEC, REFR }; // material types, used in radiance()

struct Sphere {
    double radi; // radius
    Vec p, n, e; // position, emission, color
    Refl_t refl; // reflection type (DIFFuse, SPECular, REFRactive)
    ... methods ...
    double intersect(const Ray &r) const { // returns distance, 0 if miss
    };
    double sphereHit = 0; //Name: radius, position, emission, color, material
    Sphere(int i, Vec(100,0,0),Vec(0.5,Vec(1,2,3),DIFF),0,0,0); //left
    ... initialization ...
};
```

1. S
2. Has geometric primitives: vectors, spheres, materials
3. Entirely numeric-based, no real “data structures” to speak of

└─What is smallpt anyway?



1. **S**
2. Most of the compute cost is spent in the function that traces rays.
3. is called radiance

└─What is smallpt anyway?



1. **S**
2. radiance is the function that performs this path tracing
3. Recursively calls itself a bunch of times

—What is smallpt anyway?

What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
```

```

if (
  if (
    ) if (
  )
) alias

) alias if (
    radianse
)
    radianse

if (
    )
    radianse

    radianse
    radianse
    radianse
    radianse
)

```

1. **S**
2. Recursion is guarded by a lot of control flow

What is smallpt anyway?



1. **S**
2. The control flow and computation is very numeric in nature

—What is smallpt anyway?

```

What is smallest anyway?

var radiance=0;
for (i=0; i<ray.depth; i++) {
    if (i%2==0) {
        // even depth
        // choose left child
        // ...
    } else {
        // odd depth
        // choose right child
        // ...
    }
}
return radiance;

```

1. **S**
2. uses the function `erand48` for randomness

└─What is smallpt anyway?

[illegible]

1. **S**
2. The full code continues to be more of the same

## Initial Haskell Code: radiance (1×)



1. S
2. The same computation, this time in haskell



## Initial Haskell Code: Data structures

Initial Haskell Code: Data structures

```

data Vec = Vec {-# UNPACK #-} !Double {-# UNPACK #-} !Double {-# UNPACK #-} !Double

norm :: Vec -> Vec -> Vec
[-1] :: Vec -> Double -> Vec
radial 2 * ^

lim :: Vec -> Double
norm :: Vec -> Vec
norm v = v * 1 / mag (lim v)
dot :: Vec -> Vec -> Double
mag :: Vec -> Double

data Ray = Ray Vec Vec -- origin, direction
data Refl = DIFF | SPEC | REFR -- material type, used in radiance
-- | radius, position, emission, color, refraction
data Sphere = Sphere Double Vec Vec Vec Refl

```

1. S
2. We implement the same geometric data structures in Haskell
3. Code here has an inconsistency
4. while Vec has unpack, Ray, Sphere not having unpack

Initial Haskell Code: scene data

```

spheres = [Spheres
            (
                [Spheres 1a5 (Two 1a5+1) 62.8 81.6) 0 (Two 0.75 0.25 0.25) 219F -Le6f
                [Spheres 1a5 (Two 09+1a5) 40.8 81.6) 0 (Two 0.25 0.25 0.75) 219F -Agi
                [Spheres 1a5 (Two 50 50.8 1a5) 0 0.75 219F -Dauk
                [Spheres 1a5 (Two 50 50.8 1219 1a5) 0 0.75 219F -Dauk
                [Spheres 1a5 (Two 50 50.8 81.6) 0 0.75 219F -Dauk
                [Spheres 1a5 (Two 50 50.6 1a5) 81.6) 0 0.75 219F -Dauk
                [Spheres 16.8 (Two 27 26.8 17) 0 0.999 999C -offe
                [Spheres 16.8 (Two 23 16.8 17) 0 0.999 999C -d1as
                [Spheres 602 (Two 30 602 30 81.6) 12 0 219F -2a1a
            ]
        ]
    
```

1. **S**
2. this list will be walked many times, as it contains our scene information.

## Initial Haskell code: Sphere intersection

Initial Haskell code: Sphere intersection

```

intersect :: Ray -> Sphere -> Maybe Double
intersect (Ray a d) (Sphere r p _a _c _refl) =
  if det < 0 then Nothing else f (b-sqrt(d)) (b+sqrt(d))
  where op = p - a -- Vector
        b = dot op d
        det = b*b - dot op op - r*r -- Scalar
        sdet = sqrt det
intersect r ray -> (Maybe Double, Sphere)
intersect r ray = (h, s)
  where (h,s) = foldl' f (Nothing,undefined) spheres -- Spheres iterated over
        f (h',sp) a' = case (h',intersect ray a') of
          (Nothing,Just a) -> (Just h,a')
          (Just h',Just a) | a < p -> (Just h,a')
          _ -> (h',sp)

```

1. S
2. Responsible for figuring out what the ray hits.
3. We iterate over the list of spheres.
4. Once again, numeric heavy.
5. Use a Maybe to indicate whether we've found an answer or not.

## Initial Haskell Code: radiance (1x)

[illegible]

1. **S**
2. Branch heavy
3. Recursive
4. Uses an RNG

## L

[illegible]

1. **S**
2. Use mutability to store the pixels of the image in `c`
3. Loops over all pixels in an image and shoots rays
4. Shoots `samps` number of rays per pixel and adds up the results
5. Finally, writes resulting color out by mutating the array `c`

## Initial Haskell Code: RNG (1×)

```
Initial Haskell Code: RNG (1×)  
  
foreign import ccall unsafe "erand48"  
erand48 :: Ptr CDouble -> IO Double
```

1. **S**
2. As mentioned previously, we use the RNG to decide randomly in which direction to send rays
3. erand48 is imported as a foreign ccall for parity with the C code

└ Restrict export list to main (1.13x)



1. **S**
2. The very first thing to do is to let the compiler actually optimize.
3. If a function is public, then compiler doesn't know all call sites
4. Export only the one function that's called from the outside: main
5. Allows compiler to know that other functions in module are not called from outside
6. compiler has "perfect knowledge" about these functions now

└ Restrict export list to main (1.13x)



1. **S**
2. The very first thing to do is to let the compiler actually optimize.
3. If a function is public, then compiler doesn't know all call sites
4. Export only the one function that's called from the outside: main
5. Allows compiler to know that other functions in module are not called from outside
6. compiler has "perfect knowledge" about these functions now



## Mark entries of Ray and Sphere as UNPACK and Strict (1.07x)

```
data Vec = Vec {# UNPACK a: } | Double
  (# UNPACK a: ) | Double
  (# UNPACK a: ) | Double

data Ray = Ray Vec Vec -- origin, direction
data Ray = Ray Vec Vec -- origin, direction

data Sph = SPH { SPC | SPS -- material types, used in radiance

-- radius, position, emission, color, and function
data Sphere = Sphere Double Vec Vec Vec Sph
data Sphere = Sphere {# UNPACK a: } | Double
+ {# UNPACK a: } | Vec
+ {# UNPACK a: } | Vec
+ {# UNPACK a: } | Vec Sph
+ {# UNPACK a: } | Vec Sph

strict Vec { Double a, p, m }
strict Ray { vec: Function Vec() -> v; dir: Function Vec() -> v }
strict RaySph { Double m, p, r, m;
  Double m, p, m; }
```

1. D
2. Strictness in the arguments means that they're evaluated when instantiated, not when demanded.
3. Where as Unpacking removes indirection from doing a memory lookup for components.
4. Means we have to copy everything into the data structure that it is unpacked into.
5. We don't unpack ray (Lots of calculations on its components, want those to fuse)
6. Unpack Sphere - its static from compile time
7. Don't unpack Ray, because each Vec undergoes a lot of computation.

2020-11-05

└ Use a pattern synonym to unpack Refl in Sphere (1.07x)

```

+(<# LAMBDA2D PatternSynonym R>)
data Refl = DIFF | SPEC | REFR -- material types, used in radiance
renewtype Refl = Refl Int -- material types, used in radiance
pattern DIFF, SPEC, REFR :: Refl
pattern DIFF = Refl 0
pattern SPEC = Refl 1
pattern REFR = Refl 2
+(<# COMPLETE DIFF, SPEC, REFR R>)

-- radii, position, emission, color, refraction
data Sphere = Sphere {# DIFFAC R>} Double
    {# DIFFAC R>} Vec
    {# DIFFAC R>} Vec
    {# DIFFAC R>} Vec Vec
+
+    {# DIFFAC R>} Vec Vec
+    {# DIFFAC R>} Vec {# DIFFAC R>} Vec

```

1. D
2. Was unable to unpack Refl
3. UnboxedSums are recent
4. UnboxedSums are very unpleasant
5. We're using an older trick to fake the unboxing here instead.
6. In this case it isn't much of a win, but it illustrates the technique.

└ Change from maximum on a list to max (1.08x)

Change from maximum on a list to max (1.08x)

```

--max (Vec a b c) = maximum [a,b,c]
--max (Vec a b c) = max a (max b c)

let x = a "addr" (d "maddr" t)
a = max b a "addr" p
all = if a "addr" d < 0 then a else a "maddr" (-t)
--
pt = max c
depth = depth + 1
continue f = case refi of
  STOP -> do
...
if depth > 0
then do
  w <- randomM 0 1
  let tpt = max a

```

1. D
2. Prebuild comparison
3. Don't go via list
4. GHC does not evaluate at compile time, only has RULES
5. Doesn't really help much in this case

- └ Convert erand48 to pure Haskell (1.09x)

Convert erand48 to pure Haskell (1.09x)

[illegible]

1. **S**
2. The entire premise of this talk is that Haskell can be as fast as C.
3. We're opening the black box of what `erand48` does to GHC
4. Further any impedance mismatch, such as FFI almost universally has to have, carries some bookkeeping overhead.
5. If our Haskell code was as fast as the C code moving the code into Haskell would be a win, if it was slightly slower it could still be a win.
6. Often considering your Haskell code's performance is a better option and easier than reimplementing something in C.
7. As is the way with optimizations, this is not universally true.

1. **D**
2. All these mutability locations throw in extra RTS code, extra sequencing that blocks the compiler's optimization, and dependency chains.
3. Sometimes we need mutability for performance.
4. SSA is normal to compilers though.
5. We almost start at SSA as a functional language.
6. don't break it when you don't have a good reason.

## Removing mutation: eliminate IOREf and Data.Vector.Mutable

Removing mutation: eliminate IORef and Data.Vector.Mutable

```

+ w <- VM.unsafeRead (w + h) 0
+ let c = modifiedRef 0
+ for sample [0..n-1] $ \p -> do
+   modifyRef w 0
+   for [0..m-1] $ \q -> do
+     let s = (p+1) * w + q
+     for [0..1] $ \r -> do
+       for [0..1] $ \b -> do
+         let i = (2*pi) * VM.unsafeRead w
+         img = VM.unsafeWrite (b*15+(p*20+q)) $ \p -> modifyRefReadRef p do
+           for [0..n-1] $ \x -> do
+             (rgb) = foldM0 $ \ (rgb, mut) i mp <- [0..1] last (rgb, mut) -> do
+               for r r rgb <- VM -> foldM0 $ \ (r, sample) i -> do
+                 let s = (p+1) * w + q
+                 modifyRefReadRef s
+
+...
+   modifyRefReadRef s (r, mut) -> modify (fromIntegral sample)
+   let c = VM.unsafeReadRef s 0
+   for r r rgb <- VM.unsafeReadRef s
+   VM.unsafeWriteRef s 1 $ s1 + from (clamp s1) (clamp rgb) (clamp s1) -> 0.25
+   pure (r, c + mut, s, sample) (fromIntegral sample)
+
+   pure (s1 + from (clamp s1) (clamp rgb) (clamp s1) -> 0.25)
+
+...

```

### 1. D

- Set **everything** in smallpt to be strict (1.17x)

Set everything in `smallpt` to be strict (1.17x)

[illegible]

Don't senselessly bang everything in sight.

1. **D**
2. This is not a recommendation, this is a warning.
3. We get a speedup here but it can also regress performance. Some of these bangs are regressions that are hidden.

2020-11-05

## Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
let fooOpt = \y -> y
```

1. **S**
2. Consider the function `foo` and `fooOpt`. These are equivalent
3. The fact that `x` is not used allows us to eliminate computing `x`
4. Consider the next version
5. Illegal, we need to have `x`, because it doesn't produce `ERR`
6. we can't equationally reason about the program anymore.
7. Makes it harder for `GHC`. `GHC` is conservative about bangs
8. Inhibits compiler from optimizing



## Why strictness may be bad

Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
let fooOpt = \y -> y
let foo' = let !y = error "ERR" in \y -> y
```

1. **S**
2. Consider the function `foo` and `fooOpt`. These are equivalent
3. The fact that `x` is not used allows us to eliminate computing `x`
4. Consider the next version
5. Illegal, we need to have `x`, because it doesn't produce `ERR`
6. we can't equationally reason about the program anymore.
7. Makes it harder for `GHC`. `GHC` is conservative about bangs
8. Inhibits compiler from optimizing

2020-11-05

## Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
let fooOpt = \y -> y
let foo' = let !y = error "ERR" in \y -> y
let foo'Opt = \y -> y -- ERROR! forcing foo'/'foo'Opt should give "ERR"
```

1. **S**
2. Consider the function `foo` and `fooOpt`. These are equivalent
3. The fact that `x` is not used allows us to eliminate computing `x`
4. Consider the next version
5. Illegal, we need to have `x`, because it doesn't produce `ERR`
6. we can't equationally reason about the program anymore.
7. Makes it harder for `GHC`. `GHC` is conservative about bangs
8. Inhibits compiler from optimizing

└ Reduce to only useful strictnesses in smallpt(1.17x)

```

Reduce to only useful strictnesses in smallpt(1.17x)

- let timage = sample "dir" 4
- timg = Vec S2 S2 286.6
- dir1 = norm $ Vec 0 (-0.042612) (-1)
- tca = Vec (fromIntegral u + 0.515) / (fromIntegral h) 0 0
- tcy = norm (ca "area" dir) + 0.515
-
+ let sample = sample "dir" 4
+ org = Vec S2 S2 286.6
+ dir = norm $ Vec 0 (-0.042612) (-1)
+ ca = Vec (fromIntegral u + 0.515) / (fromIntegral h) 0 0
+ cy = norm (ca "area" dir) + 0.515
+
+--
+--   r1 <- (2*) <#> random
+--   r2 <- (2*) <#> random
+--
+--   r1 <- (2*) <#> random
+--   r2 <- (2*) <#> random
+--
+--   rand <- radiance Ray (arg+1+40) (norm d1) 0
+--   rad <- radiance Ray (arg+1+40) (norm d1) 0
+--
+--   pure $! r + rad + recip (fromIntegral sample)
+--   pure $! r + Vec (clamp r1) (clamp r2) (clamp r3) + 0.25
+--   pure (r + rad + recip (fromIntegral sample))
+--   pure (r1 + Vec (clamp r1) (clamp r2) (clamp r3) + 0.25)

```

1. D
2. Thus the compiler can no longer move the computation around or simplify it.
3. Force useless work.
4. A little thinking about how the variables are used or looking at core allows us to select which ones we bang selectively.

## Use strictness strategically in entire project

```

...
- if desc() when Rayking when f (bushes) (brakes)
- where up = p - a
- up = -1e-5
-
- b = desc-up-d
- det = det-d - det-up + up
-
- f = a - if a < 0 then desc + a else if a < 0 then desc + a else Rayking
-
+ if desc()
+ when Rayking
+ when
- det fup = -1e-5
+ fup = -1e-5
+ fup = -1e-5
+ fup = -1e-5
+
- in if a < 0 then desc + a else if a < 0 then desc + a else Rayking
...

```

1. D
2. Sometimes (point out 'intersect') we have to rearrange the code though when we use bangs.
3. Bangs tell the compiler to make more efficient code, but take away the compiler's options in how to do so.
4. Only take away the compiler's liberties when it's using them poorly.
5. Becomes intuitive.

- Remove Maybe from intersect(s) ( $1.32\times$ )

[illegible]

1. **S**
2. This is a far more performance critical version of what we saw with 'maximum' vs. 'max'.
3. innermost functions are of critical importance. remove Maybe which significantly reduces the boxing
4. Since a Ray that fails to intersect something can be said to intersect at infinity, Double already actually covers the structure at play
5. This also reduces allocation.

Hand unroll the fold in intersects ( $1.35\times$ )

[illegible]

1. **D**
2. 'intersects' is very hot
3. Loop unrolling
4. Many compilers do this for us, and there are special versions of it like Duff's Device.
5. Sadly GHC doesn't
6. Can do variants by hand.
7. RULE could handle each one specifically (only exactly that one)?

## Custom datatype for intersects parameter passing

Custom datatype for intersects parameter passing

Std: Tuple with possibly-unevaluated Double and Sphere  
 Req: Reference to a guaranteed-to-be-evaluated Double and Sphere  
 intersects :: Ray -> (Double, Sphere)  
 -data T = T (Double) (Sphere)  
 +  
 +intersects :: Ray -> T  
 intersects ray =  
 + if [...] then (intersect ray sphLeft, sphLeft) sphRight) ... sphLite  
 + if [...] then (intersect ray sphLeft) sphRight) sphLite  
 + else  
 + if (h', sp) s' =  
 + let ts = intersect ray s' in if x < h' then (x, s') else (h', sp)  
 + if (h' s' sp) ts =  
 + let ts = intersect ray s' in if x < h' then T x s' else T h' sp  
 +  
 +radiuses :: Ray -> Int -> Random Vec  
 radiuses rpfRay n d1 depth = case intersects ray of  
 + (ts,\_) | t == 1/0.0 -> return 0  
 + (ts, (Sphere \_ p s rad)) -> do  
 + if t < .3 | t == 1/0.0 -> return 0  
 + if t (Sphere \_ p s rad) -> do  
 + let ts = s - d \* s  
 + ts = norm 0 s - p  
 + ts = if dot s d < 0 then s else negate s

1. D
2. We can optimize data passing.
3. Want: Data strict, but not unpacked.
4. Compiler knows its evaluated but no copying
5. A normal tuple lacks strictness information.
6. An unboxed tuple forces copying
7. Strict Tuple.
8. This exists in libraries of course, but we wanted to illustrate it.

## └ Optimize file writing

Optimize file writing

```
build-dependencies:
  base >= 4.12 && < 4.15
+ , bytestring >= 0.11

-letat :: Double -> Int
-letat x = floor $ clamp x ** recip 2.2 * 255 + 0.5
+letat :: Double -> BB.Builder -- O(1) concatenation
+letat x = BB.intDec (floor (clamp x ** recip 2.2 * 255 + 0.5)) <> BB.char8 ' '
+...

withFile "image.ppm" WriteMode $ \hdl -> do
-  hPutStrLn hdl "P3\n#RGB\n"
-  for_ img \Vec x g b -> do
-    hPrintf hdl "%d %d %d\n" (toNat x) (toNat g) (toNat b)
-    BB.BuilderBuilder hdl $
+  BB.stringBuilder hdl $
+    BB.intDec x <> BB.char8 ' ' <> BB.intDec g <> BB.char8 'a' <>
+    BB.intDec 255 <> BB.char8 '\n' <>
+    (concat $ map (\Vec x g b -> letat x <> letat g <> letat b) img)
```

1. D
2. Strings are inefficient
3. 'bytestring' has some efficient writing code, so we just convert to that for a modest gain.



└ Use LLVM backend (1.87x)


Use LLVM backend (1.87x)

```
package smallpt-opt  
+ ghc-options: -fllvm
```

1. **S**
2. Finally, this particular code is quite numeric heavy.
3. There are optimizations for numeric heavy code we're missing in GHC.
4. LLVM has an extensive library of laws to optimize low level numeric ops.
5. LLVM is too low-level to understand haskell as haskell.
6. LLVM makes decisions with the tacit assumption that the assembly came from a C-like language, which is often to the detriment of a Haskell-like language.
7. In this case, as the code is "fortran-like", LLVM wins.

## The view from the mountaintop

### 1. D



└─ Avoid CPU ieee754 slow paths

Avoid CPU ieee754 slow paths

```
intersect :: Ray -> Sphere -> Double
intersect (Ray u d) (Sphere s p _c _rad1) =
  if succ0
  - then 1/0.0
  + else 1e20
  else
    ...
    - in if s>eps then a else if s>eps then a else 1/0.0
    + in if s>eps then a else if s>eps then a else 1e20
  ...

radiance :: Ray -> Int -> RandomM Vec
radiance ray@(Ray u d) depth = case intersects ray of
- (f t _) | t == 1/0.0 -> return 0
+ (f 1e20 _) -> return 0
  ...
```

1. D
2. We used +Inf to match the Maybeness
3. C++ code set 1e20 s the horizon
4. Mechanical sympathy is important.
5. Know how the CPU (abstractly) executes - slow path / fast path.

## └ Fix differences with C++ version

Fix differences with C++ version

```
- if depth>2  
+ if depth>2 -- depth = depth + 1  
...
```

1. **S**
2. Since the sha1 of the output didn't match the C++ version we started investigating.
3. clang++, g++ actually produce different sha1s
4. unincremented depth was being used in one branch, causing us to do more work
5. now confident to say we're doing the same computation as C++

## Takeaways

Takeaways	<ul style="list-style-type: none"><li>▶ The wording in 'Interacts' is ugly.</li><li>▶ (We feel) the maintainability of this code hasn't been significantly harmed.</li><li>▶ We're faster than clang++ and within 6% of g++.</li><li>▶ Haven't exhausted the optimization opportunities.</li><li>▶ GHC could have to do <i>several</i> of these optimizations for us.</li><li>▶ Others are just good Haskell style.</li><li>▶ Clean Haskell is often performant Haskell.</li><li>▶ Repository stepping through each optimization is available at ...</li></ul>
-----------	--

### 1. D