

Optimizing smallpt


Davean Scies, Siddharth Bhat

Haskell Exchange


November 4th, 2020

What is smallpt anyway?

What is smallpt anyway?




What is smallpt anyway?



- 100 LoC C demo of a raytracer

What is smallpt anyway?



- 100 LoC C demo of a raytracer
- Perfect for an optimization case study

What is smallpt anyway?

```

#include <math.h>
#include <stdlib.h>
#include <stdio.h>
struct Vec {
    double x, y, z; // position, also color (r,g,b)
    ... methods...
};
struct Ray { Vec o, d; Ray(Vec o_, Vec d_) : o(o_), d(d_) {} };
enum Refl_t { DIFF, SPEC, REFR }; // material types, used in radiance()
struct Sphere {
    double rad; // radius
    Vec p, e, c; // position, emission, color
    Refl_t refl; // reflection type (DIFFuse, SPECular, REFRactive)
    ... methods ...
    double intersect(const Ray &r) const // returns distance, 0 if nohit
};
Sphere spheres[] = { //Scene: radius, position, emission, color, material
    Sphere(1e5, Vec( 1e5+1,40.8,81.6), Vec(),Vec(.75,.25,.25),DIFF), //Left
    ... initialization ...
};
inline bool intersect(const Ray &r, double &t, int &id)

```

What is smallpt anyway?

```

Vec radiance(const Ray &r, int depth, unsigned short *Xi){
    double t;                                // distance to intersection
    int id=0;                                // id of intersected object
    if (!intersect(r, t, id)) return Vec(); // if miss, return black
    const Sphere &obj = spheres[id];         // the hit object
    Vec x=r.o+r.d*t, n=(x-obj.p).norm(), nl=n.dot(r.d)<0?n:n*-1, f=obj.c;
    double p = f.x>f.y && f.x>f.z ? f.x : f.y>f.z ? f.y : f.z; // max refl
    if (++depth>5) if (erand48(Xi)<p) f=f*(1/p); else return obj.e; //R.R.
    if (obj.refl == DIFF){                      // Ideal DIFFUSE reflection
        double r1=2*M_PI*erand48(Xi), r2=erand48(Xi), r2s=sqrt(r2);
        Vec w=nl, u=((fabs(w.x)>.1?Vec(0,1):Vec(1))%w).norm(), v=w%u;
        Vec d = (u*cos(r1)*r2s + v*sin(r1)*r2s + w*sqrt(1-r2)).norm();
        return obj.e + f.mult(radiance(Ray(x,d),depth,Xi));
    } else if (obj.refl == SPEC)                // Ideal SPECULAR reflection
        return obj.e + f.mult(radiance(Ray(x,r.d-n*2*n.dot(r.d)),depth,Xi));
    Ray reflRay(x, r.d-n*2*n.dot(r.d));        // Ideal dielectric REFRACTION
    bool into = n.dot(nl)>0;                    // Ray from outside going in?
    double nc=1, nt=1.5, nnt=into?nc/nt:nt/nc, ddn=r.d.dot(nl), cos2t;
    if ((cos2t=1-nnt*nnt*(1-ddn*ddn))<0)         // Total internal reflection
        return obj.e + f.mult(radiance(reflRay,depth,Xi));
    Vec tdir = (r.d*nnt - n*((into?-1:1)*(ddn*nnt+sqrt(cos2t)))).norm();
    double a=nt-nc, b=nt+nc, R0=a*a/(b*b), c = 1-(into?-ddn:tdir.dot(n));
    double Re=R0+(1-R0)*c*c*c*c*c,Tr=1-Re,P=.25+.5*Re,RP=Re/P,TP=Tr/(1-P);
    return obj.e + f.mult(depth>2 ? (erand48(Xi)<P ? // Russian roulette
        radiance(reflRay,depth,Xi)*RP:radiance(Ray(x,tdir),depth,Xi)*TP) :
        radiance(reflRay,depth,Xi)*Re+radiance(Ray(x,tdir),depth,Xi)*Tr);
}

```

Establishing baselines

```
main :: IO ()
main = smallptM 200 200 256

smallptM :: Int -> Int -> Int -> IO ()
smallptM !w !h !nsamps = do
  -- ... processing
  withFile "image.ppm" WriteMode $ \hdl -> do
    hPrintf hdl "P3\n%d %d\n%d\n" w h (255::Int)
    flip mapM_ [0..w*h-1] $ \i -> do
      Vec r g b <- VM.unsafeRead c i
      hPrintf hdl "%d %d %d " (toInt r) (toInt g) (toInt b)
```


Haskell: the first stab

```

radiance :: Ray -> CInt -> Ptr CUShort -> IO Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
  (Nothing, _) -> return zeroV
  (Just t, Sphere _r p e c refl) -> do
    let x = o `addv` (d `mulvs` t)
        n = norm $ x `subv` p
        nl = if n `dot` d < 0 then n else n `mulvs` (-1)
        pr = maxv c
    depth' = depth + 1
    continue f = case refl of
      DIFF -> do
        r1 <- ((2*pi)* ) `fmap` erand48 xi
        r2 <- erand48 xi
        let r2s = sqrt r2
            w@(Vec wx _ _) = nl
            u = norm $ (if abs wx > 0.1 then (Vec 0 1 0) else (Vec 1 0 0))
            v = w `cross` u
            d' = norm $ (u `mulvs` (cos r1*r2s)) `addv` (v `mulvs` (sin r1*r2s))
        rad <- radiance (Ray x d') depth' xi
        return $ e `addv` (f `mulv` rad)

```

SPEC -> do

```

let d' = d `subv` (n `mulvs` (2 * (n `dot` d)))

```

Restrict export list to 'main'

```
-module Main where  
+module Main (main) where
```

Mark entries of Ray and Sphere as UNPACK and Strict.

```
-data Ray = Ray Vec Vec -- origin, direction
+data Ray = Ray {-# UNPACK #-} !Vec {-# UNPACK #-} !Vec -- origin, direction

data Refl = DIFF | SPEC | REFR -- material types, used in radiance

-- radius, position, emission, color, reflection
-data Sphere = Sphere Double Vec Vec Vec !Refl
+data Sphere = Sphere {-# UNPACK #-} !Double {-# UNPACK #-} !Vec {-# UNPACK #-} !Vec
```

Use a pattern synonym to unpack Refl in Sphere.

```

+{-# LANGUAGE PatternSynonyms #-}

% e77b26f
-data Refl = DIFF | SPEC | REFR -- material types, used in radiance
+newtype Refl = Refl Int -- material types, used in radiance
+pattern DIFF,SPEC,REFR :: Refl
+pattern DIFF = Refl 0
+pattern SPEC = Refl 1
+pattern REFR = Refl 2
+{-# COMPLETE DIFF, SPEC, REFR #-}

-- radius, position, emission, color, reflection
-data Sphere = Sphere {-# UNPACK #-} !Double {-# UNPACK #-} !Vec {-# UNPACK #-}
+data Sphere = Sphere {-# UNPACK #-} !Double {-# UNPACK #-} !Vec {-# UNPACK #-}

intersect :: Ray -> Sphere -> Maybe Double
intersect (Ray o d) (Sphere r p _e _c _refl) =

```

Change from maximum on a list to max

```

-maxv (Vec a b c) = maximum [a,b,c]
+maxv (Vec a b c) = max a (max b c)

@@ -84,7 +85,6 @@ radiance ray@(Ray o d) depth xi = case intersects ray of
    let x = o `addv` (d `mulvs` t)
        n = norm $ x `subv` p
        nl = if n `dot` d < 0 then n else n `mulvs` (-1)
-    pr = maxv c
    depth' = depth + 1
    continue f = case refl of
        DIFF -> do
@@ -140,6 +140,7 @@ radiance ray@(Ray o d) depth xi = case intersects ray of
    if depth'>5
    then do
        er <- erand48 xi
+    let !pr = maxv c

```

Convert erand48 to pure Haskell

```

-radiance :: Ray -> CInt -> Ptr CUSHort -> IO Vec
+radiance :: Ray -> Int -> IORef Word64 -> IO Vec
  radiance ray@(Ray o d) depth xi = case intersects ray of
    (Nothing,_) -> return zeroV
    (Just t,Sphere _r p e c refl) -> do
@@ -153,9 +153,8 @@ smallpt w h nsamps = do
    cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
    cy = norm (cx `cross` dir) `mulvs` 0.5135
    c <- VM.replicate (w * h) zeroV
-   allocaArray 3 $ \xi ->
-     flip mapM_ [0..h-1] $ \y -> do
-       --hPrintf stderr "\rRendering (%d spp) %5.2f%%" (samps*4::Int) (100.0*
+   xi <- newIORef 0
+   flip mapM_ [0..h-1] $ \y -> do
    writeXi xi y

```

TODO: add better version of erand48 into this commit

Change erand48 to IORefU.

```

-radiance :: Ray -> Int -> IORef Word64 -> IO Vec
+radiance :: Ray -> Int -> IORefU Word64 -> IO Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
    (Nothing,_) -> return zeroV
    (Just t,Sphere _r p e c refl) -> do
@@ -153,7 +154,7 @@ smallpt w h nsamps = do
    cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
    cy = norm (cx `cross` dir) `mulvs` 0.5135
    c <- VM.replicate (w * h) zeroV
- xi <- newIORef 0
+ xi <- newIORefU 0
    flip mapM_ [0..h-1] $ \y -> do
        writeXi xi y
        flip mapM_ [0..w-1] $ \x -> do
@@ -181,8 +182,8 @@ smallpt w h nsamps = do
        Vec r g b <- VM.unsafeRead c i
        hPrintf hdl "%d %d %d " (toInt r) (toInt g) (toInt b)

-writeXi :: IORef Word64 -> Int -> IO ()
-writeXi !xi !y = writeIORef xi (mkErand48Seed' y)
+writeXi :: IORefU Word64 -> Int -> IO ()
+writeXi !xi !y = writeIORefU xi (mkErand48Seed' y)

```

Rewrite the remaining IORef into a foldM

TODO

Remove the Data.Vector.Mutable by being purer

```
-      r <- newIORef zerov
-      flip mapM_ [0..samps-1] $ \_s -> do
+      Vec rr rg rb <- (\f -> foldM f zerov [0..samps-1]) $ \ !r _s -> do
          r1 <- (2*) `fmap` erand48 xi
          let dx = if r1<1 then sqrt r1-1 else 1-sqrt(2-r1)
          r2 <- (2*) `fmap` erand48 xi
@@ -171,9 +170,8 @@ smallpt w h nsamps = do
            (cy `mulvs` (((sy + 0.5 + dy)/2 + fromIntegral y)/fromIntegral w))
            rad <- radiance (Ray (org`addv`(d`mulvs`140)) (norm d)) 0 xi
            -- Camera rays are pushed forward ~~~~~ to start in interior
-      modifyIORef r (\adv -> (rad `mulvs` (1 / fromIntegral samps)))
+      pure (r `addv` (rad `mulvs` (1 / fromIntegral samps)))
-      ci <- VM.unsafeRead c i
-      Vec rr rg rb <- readIORef r
-      VM.unsafeWrite c i $ ci `addv` (Vec (clamp rr) (clamp rg) (clamp r
```

Set everything in smallpt to be strict

```

- let samps = nsamps `div` 4
- org = Vec 50 52 295.6
- dir = norm $ Vec 0 (-0.042612) (-1)
- cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
- cy = norm (cx `cross` dir) `mulvs` 0.5135
+ let !samps = nsamps `div` 4
+ !org = Vec 50 52 295.6
+ !dir = norm $ Vec 0 (-0.042612) (-1)
+ !cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
+ !cy = norm (cx `cross` dir) `mulvs` 0.5135
-
- r1 <- (2*) `fmap` erand48 xi
- let dx = if r1<1 then sqrt r1-1 else 1-sqrt(2-r1)
- r2 <- (2*) `fmap` erand48 xi
- let dy = if r2<1 then sqrt r2-1 else 1-sqrt(2-r2)
- d = (cx `mulvs` (((sx + 0.5 + dx)/2 + fromIntegral x)/fromIntegral w
-      (cy `mulvs` (((sy + 0.5 + dy)/2 + fromIntegral y)/fromIntegral h)
-      rad <- radiance (Ray (org`addv` (d`mulvs`140)) (norm d)) 0 xi
+ !r1 <- (2*) `fmap` erand48 xi
+ let !dx = if r1<1 then sqrt r1-1 else 1-sqrt(2-r1)
+ !r2 <- (2*) `fmap` erand48 xi
+ let !dy = if r2<1 then sqrt r2-1 else 1-sqrt(2-r2)
+ !d = (cx `mulvs` (((sx + 0.5 + dx)/2 + fromIntegral x)/fromIntegral w
+      (cy `mulvs` (((sy + 0.5 + dy)/2 + fromIntegral y)/fromIntegral h)

```

Reduce to only effectful strictnesses

```

- let !samps = nsamps `div` 4
-   !org = Vec 50 52 295.6
-   !dir = norm $ Vec 0 (-0.042612) (-1)
-   !cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
-   !cy = norm (cx `cross` dir) `mulvs` 0.5135
+ let samps = nsamps `div` 4
+   org = Vec 50 52 295.6
+   dir = norm $ Vec 0 (-0.042612) (-1)
+   cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
+   cy = norm (cx `cross` dir) `mulvs` 0.5135
-
-       !r1 <- (2*) `fmap` erand48 xi
+       r1 <- (2*) `fmap` erand48 xi
-
-       !r2 <- (2*) `fmap` erand48 xi
+       r2 <- (2*) `fmap` erand48 xi
-
-       !rad <- radiance (Ray (org`addv`(d`mulvs`140)) (norm d)) 0 xi
+       rad <- radiance (Ray (org`addv`(d`mulvs`140)) (norm d)) 0 xi
-
-       pure $! (r `addv` (rad `mulvs` (1 / fromIntegral samps)))
-       pure $! ci `addv` (Vec (clamp rr) (clamp rg) (clamp rb) `mulvs` 0.
+       pure $ (r `addv` (rad `mulvs` (1 / fromIntegral samps)))
+       pure $ ci `addv` (Vec (clamp rr) (clamp rg) (clamp rb) `mulvs` 0.2

```

Remove Maybe from intersect(s)

```

-intersect :: Ray -> Sphere -> Maybe Double
-intersect (Ray o d) (Sphere r p _e _c _refl) =
-  if det<0 then Nothing else f (b-sdet) (b+sdet)
-  where op = p `subv` o
-         eps = 1e-4
-         b = op `dot` d
-         det = b*b - (op `dot` op) + r*r
-         sdet = sqrt det
-         f a s = if a>eps then Just a else if s>eps then Just s else Nothing
-
+intersect :: Ray -> Sphere -> Double
+intersect (Ray o d) (Sphere r p _e _c _refl) =
+  if det<0 then (1/0.0) else f (b-sdet) (b+sdet)
+  where op = p `subv` o
+         eps = 1e-4
+         b = op `dot` d
+         det = b*b - (op `dot` op) + r*r
+         sdet = sqrt det
+         f a s = if a>eps then a else if s>eps then s else (1/0.0)
+
-intersects :: Ray -> (Maybe Double, Sphere)
+intersects :: Ray -> (Double, Sphere)

```

Hand unroll the fold in intersects

```

intersects :: Ray -> (Double, Sphere)
-intersects ray = (k, s)
-  where (k,s) = foldl1' f (1/0.0,undefined) spheres
-          f (k', sp) s' = let !x = intersect ray s' in if x < k' then (x, s') el
+intersects ray =
+  f (f (f (f (f (f (f (f (intersect ray sphLeft, sphLeft) sphRight) sphBack)
+  where
+  f (k', sp) s' = let !x = intersect ray s' in if x < k' then (x, s') else (
-spheres :: [Sphere]
-spheres = let s = Sphere ; z = zerov ; (.* ) = mulvs ; v = Vec in
-  [ s 1e5 (v (1e5+1) 40.8 81.6)      z (v 0.75 0.25 0.25) DIFF --Left
-  , s 1e5 (v (-1e5+99) 40.8 81.6)    z (v 0.25 0.25 0.75) DIFF --Right
-  , s 1e5 (v 50 40.8 1e5)            z (v 0.75 0.75 0.75) DIFF --Back
-  , s 1e5 (v 50 40.8 (-1e5+170))     z z                                DIFF --Frnt
-  , s 1e5 (v 50 1e5 81.6)            z (v 0.75 0.75 0.75) DIFF --Botm
-  , s 1e5 (v 50 (-1e5+81.6) 81.6)    z (v 0.75 0.75 0.75) DIFF --Top
-  , s 16.5(v 27 16.5 47)             z ((v 1 1 1).* 0.999) SPEC --Mirr
-  , s 16.5(v 73 16.5 78)            z ((v 1 1 1).* 0.999) REFR --Glas
-  , s 600 (v 50 (681.6-0.27) 81.6)  (v 12 12 12)          z DIFF]--Lite
+sphLeft, sphRight, sphBack, sphFrnt, sphBotm, sphTop, sphMirr, sphGlas, sphLit
+sphLeft  = Sphere 1e5  (Vec (1e5+1) 40.8 81.6)    zerov (Vec 0.75 0.25 0.25) D
+sphRight = Sphere 1e5  (Vec (-1e5+99) 40.8 81.6)  zerov (Vec 0.25 0.25 0.75) D
+sphBack  = Sphere 1e5  (Vec 50 40.8 1e5)          (Vec 0.75 0.75 0.75) D
+sphFrnt  = Sphere 1e5  (Vec 50 40.8 (-1e5+170))   z          (Vec 0.75 0.75 0.75) D
+sphBotm  = Sphere 1e5  (Vec 50 1e5 81.6)           z          (Vec 0.75 0.75 0.75) D
+sphTop    = Sphere 1e5  (Vec 50 (-1e5+81.6) 81.6)  z          (Vec 0.75 0.75 0.75) D
+sphMirr   = Sphere 16.5 (Vec 27 16.5 47)          z          (Vec 1 1 1) SPEC
+sphGlas   = Sphere 16.5 (Vec 73 16.5 78)          z          (Vec 1 1 1) REFR
+sphLit    = Sphere 600  (Vec 50 (681.6-0.27) 81.6) (Vec 12 12 12) z DIFF

```

Marking intersects' f parameters strict

```

intersects ray =
    f (f (f (f (f (f (f (intersect ray sphLeft, sphLeft) sphRight) sphBack)
    where
-   f (k', sp) s' = let !x = intersect ray s' in if x < k' then (x, s') else (
+   f (!(k'), !sp) !s' = let !x = intersect ray s' in if x < k' then (x, s') el

```


Strategic application of strictness

```

-   if det<0 then (1/0.0) else f (b-sdet) (b+sdet)
-   where op = p `subv` o
-         eps = 1e-4
-         b = op `dot` d
-         det = b*b - (op `dot` op) + r*r
-         sdet = sqrt det
-         f a s = if a>eps then a else if s>eps then s else (1/0.0)
+   if det<0
+   then (1/0.0)
+   else
+     let !eps = 1e-4
+         !sdet = sqrt det
+         !a = b-sdet
+         !s = b+sdet
+     in if a>eps then a else if s>eps then s else (1/0.0)
+   where
+     !det = b*b - (op `dot` op) + r*r
+     !b = op `dot` d
+     !op = p `subv` o
-   (t,_) | t == (1/0.0) -> return zeroV
-   (t,Sphere _r p e c refl) -> do
-     let x = o `addv` (d `mulvs` t)

```

The view from the mountaintop



Takeaways

Takeaways

- Haskell can be fast

Takeaways

- Haskell can be fast
- ... with a lot of work!

Takeaways

- Haskell can be fast
- ... with a lot of work!
- Accumulate optimizations to accrue performance wins.

Takeaways

- Haskell can be fast
- ... with a lot of work!
- Accumulate optimizations to accrue performance wins.
- [Raw Google Sheet of our transformations](#)
- github.com/bollu/smallpt-opt