

Optimizing smallpt

Davean Scies, Siddharth Bhat

Haskell Exchange

November 4th, 2020


What is smallpt anyway?

What is smallpt anyway?



- ▶ 99 LoC C++: **small** path tracer.
- ▶ Ported to many languages, including Haskell! (Thanks to Vo Minh Thu/noteed).
- ▶ Start from noteed's original source; SHA the output image from the Haskell source for baseline.

A TL;DR of a path tracer




A TL;DR of a path tracer




void ← D

black (0,0,0)


A TL;DR of a path tracer



A TL;DR of a path tracer



A TL;DR of a path tracer



- ▶ When to end recursion?
- ▶ Choose to end recursion *randomly* after some large number of rounds
- ▶ “Russian Roulette”

A TL;DR of a path tracer

- ▶ Recursive (recursively send more light rays)
- ▶ Branching (hit an object? is it a light source? what material?)
- ▶ Number crunching (reflection, refraction, sphere-ray-intersection)
- ▶ Randomness (when do we stop a light ray)

What is smallpt anyway?

```
struct Vec {
    double x, y, z; // position, also color (r,g,b)
    ... methods...
};

struct Ray { Vec o, d; Ray(Vec o_, Vec d_) : o(o_), d(d_) {} };

enum Refl_t { DIFF, SPEC, REFR }; // material types, used in radiance()

struct Sphere {
    double rad; // radius
    Vec p, e, c; // position, emission, color
    Refl_t refl; // reflection type (DIFFuse, SPECular, REFRactive)
    ... methods ...
    double intersect(const Ray &r) const // returns distance, 0 if nohit
};

Sphere spheres[] = {//Scene: radius, position, emission, color, material
    Sphere(1e5, Vec( 1e5+1,40.8,81.6), Vec(),Vec(.75,.25,.25),DIFF),//Left
    ... initialization ...
};
```

What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
```

}

What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
```

```
    radiance
```

```
    radiance
```

```
    radiance
```

```
    radiance  
    radiance
```

```
    radiance  
    radiance
```

```
}
```

What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
```

```
    if (          ) if (          )          else  
    if (          ){
```

```
        radiance  
    } else if (          )  
        radiance
```

```
    if (          )  
        radiance
```

```
        radiance          radiance  
        radiance          radiance  
    }
```

What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
```

```
Vec x=r.o+r.d*t, n=(x-obj.p).norm(), nl=n.dot(r.d)<0?n:n*-1, f=obj.c;
```

```
if (          ) if (          )          else  
if (          ){
```

```
          radiance  
} else if (          )  
          radiance
```

```
if ((cos2t=1-nnt*nnt*(1-ddn*ddn))<0)  
    radiance
```

```
    radiance          radiance  
    radiance          radiance  
}
```

What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){

    Vec x=r.o+r.d*t, n=(x-obj.p).norm(), nl=n.dot(r.d)<0?n:n*-1, f=obj.c;

    if (          ) if (erand48(Xi) )          else
    if (          ){
        erand48(Xi)          erand48(Xi)

        radiance
    } else if (          )
        radiance

    if ((cos2t=1-nnt*nnt*(1-ddn*ddn))<0)
        radiance

        erand48(Xi)
    radiance
    radiance
    radiance
}
```


What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
    double t; // distance to intersection
    int id=0; // id of intersected object
    if (!intersect(r, t, id)) return Vec(); // if miss, return black
    const Sphere &obj = spheres[id]; // the hit object
    Vec x=r.o+r.d*t, n=(x-obj.p).norm(), nl=n.dot(r.d)<0?n:n*-1, f=obj.c;
    double p = f.x>f.y && f.x>f.z ? f.x : f.y>f.z ? f.y : f.z; // max refl
    if (++depth>5) if (erand48(Xi)<p) f=f*(1/p); else return obj.e; //R.R.
    if (obj.refl == DIFF){ // Ideal DIFFUSE reflection
        double r1=2*M_PI*erand48(Xi), r2=erand48(Xi), r2s=sqrt(r2);
        Vec w=nl, u=((fabs(w.x)>.1?Vec(0,1):Vec(1))%w).norm(), v=w%u;
        Vec d = (u*cos(r1)*r2s + v*sin(r1)*r2s + w*sqrt(1-r2)).norm();
        return obj.e + f.mult(radiance(Ray(x,d),depth,Xi));
    } else if (obj.refl == SPEC) // Ideal SPECULAR reflection
        return obj.e + f.mult(radiance(Ray(x,r.d-n*2*n.dot(r.d)),depth,Xi));
    Ray reflRay(x, r.d-n*2*n.dot(r.d)); // Ideal dielectric REFRACTION
    bool into = n.dot(nl)>0; // Ray from outside going in?
    double nc=1, nt=1.5, nnt=into?nc/nt:nt/nc, ddn=r.d.dot(nl), cos2t;
    if ((cos2t=1-nnt*nnt*(1-ddn*ddn)<0) // Total internal reflection
        return obj.e + f.mult(radiance(reflRay,depth,Xi));
    Vec tdir = (r.d*nnt - n*((into?-1):1)*(ddn*nnt+sqrt(cos2t))).norm();
    double a=nt-nc, b=nt+nc, R0=a*a/(b*b), c = 1-(into?-ddn:tdir.dot(n));
    double Re=R0+(1-R0)*c*c*c*c*c, Tr=1-Re, P=.25+.5*Re, RP=Re/P, TP=Tr/(1-P);
    return obj.e + f.mult(depth>2 ? (erand48(Xi)<P ? // Russian roulette
        radiance(reflRay,depth,Xi)*RP:radiance(Ray(x,tdir),depth,Xi)*TP) :
        radiance(reflRay,depth,Xi)*Re+radiance(Ray(x,tdir),depth,Xi)*Tr);
}
```

Initial Haskell Code: radiance (1×)

```
radiance :: Ray -> CInt -> Ptr CUSHort -> IO Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
  (Nothing,_) -> return zerov
  (Just t,Sphere _r p e c refl) -> do
```

```
  continue f = case refl of - BRANCHING
    DIFF -> do
      r1 <- ((2*pi)*) `fmap` erand48 xi -- RNG
```

```
      radiance
```

```
    SPEC -> do
```

```
      rad <- radiance -- RECURSION
```

```
    REFR -> do
```

```
      if
        then do
          rad <- radiance
        ...
```

Initial Haskell Code: Data structures (1×)

```
data Vec = Vec {-# UNPACK #-} !Double {-# UNPACK #-} !Double {-# UNPACK #-} !Double

cross :: Vec -> Vec -> Vec
(.*) :: Vec -> Double -> Vec
infixl 7 .*
len :: Vec -> Double
norm :: Vec -> Vec
norm v = v .* recip (len v)
dot :: Vec -> Vec -> Double
maxv :: Vec -> Double

data Ray = Ray Vec Vec -- origin, direction
data Refl = DIFF | SPEC | REFR -- material types, used in radiance
-- / radius, position, emission, color, reflection
data Sphere = Sphere Double Vec Vec Vec Refl
```

Initial Haskell Code: scene data (1×)

```
spheres :: [Sphere]
spheres =
  [ Sphere 1e5  (Vec (1e5+1) 40.8 81.6) 0 (Vec 0.75 0.25 0.25) DIFF --Left
  , Sphere 1e5  (Vec (99-1e5) 40.8 81.6) 0 (Vec 0.25 0.25 0.75) DIFF --Right
  , Sphere 1e5  (Vec 50 40.8 1e5)          0 0.75 DIFF --Back
  , Sphere 1e5  (Vec 50 40.8 (170-1e5))    0 0 DIFF --Frnt
  , Sphere 1e5  (Vec 50 1e5 81.6)          0 0.75 DIFF --Botm
  , Sphere 1e5  (Vec 50 (81.6-1e5) 81.6)   0 0.75 DIFF --Top
  , Sphere 16.5 (Vec 27 16.5 47)           0 0.999 SPEC --Mirr
  , Sphere 16.5 (Vec 73 16.5 78)           0 0.999 REFR --Glas
  , Sphere 600  (Vec 50 681.33 81.6)      12 0 DIFF --Lite
```

Initial Haskell code: Sphere intersection

```
intersect :: Ray -> Sphere -> Maybe Double
intersect (Ray o d) (Sphere r p _e _c _refl) =
    if det<0 then Nothing else f (b-sdet) (b+sdet)
    where op = p - o -- Numeric
          eps = 1e-4
          b = dot op d
          det = b*b - dot op op + r*r -- Numeric
          sdet = sqrt det
          f a s = if a>eps then Just a else if s>eps then Just s else Nothing
intersects :: Ray -> (Maybe Double, Sphere)
intersects ray = (k, s)
    where (k,s) = foldl' f (Nothing,undefined) spheres -- Spheres iterated over
          f (k',sp) s' = case (k',intersect ray s') of
              (Nothing,Just x) -> (Just x,s')
              (Just y,Just x) | x < y -> (Just x,s')
              _ -> (k',sp)
```

Initial Haskell Code: radiance (1×)

```
radiance :: Ray -> CInt -> Ptr CUSHort -> IO Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
  (Nothing,_) -> return zeroV
  (Just t,Sphere _r p e c refl) -> do
    let x = o `addv` (d `mulvs` t)
        n = norm $ x `subv` p
        nl = if n `dot` d < 0 then n else n `mulvs` (-1)
        pr = maxv c
    depth' = depth + 1
    continue f = case refl of
      DIFF -> do
        r1 <- ((2*pi)* `fmap` erand48 xi)
        r2 <- erand48 xi
        let r2s = sqrt r2
            w@(Vec wx _ _) = nl
            u = norm $ (if abs wx > 0.1 then (Vec 0 1 0) else (Vec 1 0 0)) `cross` w
            v = w `cross` u
            d' = norm $ (u`mulvs`(cos r1*r2s)) `addv` (v`mulvs`(sin r1*r2s)) `addv` (w`mulvs`sqrt (1-r2))
        rad <- radiance (Ray x d') depth' xi
        return $ e `addv` (f `mulv` rad)
      SPEC -> do
        let d' = d `subv` (n `mulvs` (2 * (n`dot`d)))
        rad <- radiance (Ray x d') depth' xi
        return $ e `addv` (f `mulv` rad)
      REFR -> do
        let reflRay = Ray x (d `subv` (n `mulvs` (2* n`dot`d)))
            into = n`dot`nl > 0
            nc = 1
            nt = 1.5
            nnt = if into then nc/nt else nt/nc
            ddn = d`dot`nl
            cos2t = 1-nnt*nnt*(1-ddn*ddn)
        if cos2t<0
        then do
          rad <- radiance reflRay depth' xi
        ...
```

Initial Haskell Code: Entry point (1×)

```
smallpt :: Int -> Int -> Int -> IO ()
smallpt w h nsamps = do
  ...
  c <- VM.replicate (w * h) 0
  allocArray 3 \xi -> -- Create mutable memory
  flip mapM_ [0..h-1] $ \y -> do -- Loop
    writeXi xi y
    for_ [0..w-1] \x -> do -- Loop
      let i = (h-y-1) * w + x
      for_ [0..1] \sy -> do -- Loop
        for_ [0..1] \sx -> do -- Loop
          r <- newIORef 0 -- Create mutable memory
          for_ [0..samps-1] \_s -> do -- Loops, Loops
            r1 <- (2*) <$> erand48 xi
            ...
            rad <- radiance (Ray (org+d.*140) (norm d)) 0 xi -- Crunch
            ...
            modifyIORef r (+ rad .* recip (fromIntegral samps)) -- Write
          ci <- VM.unsafeRead c i
          Vec rr rg rb <- readIORef r
          VM.unsafeWrite c i $
            ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25 -- Write
        ...
```

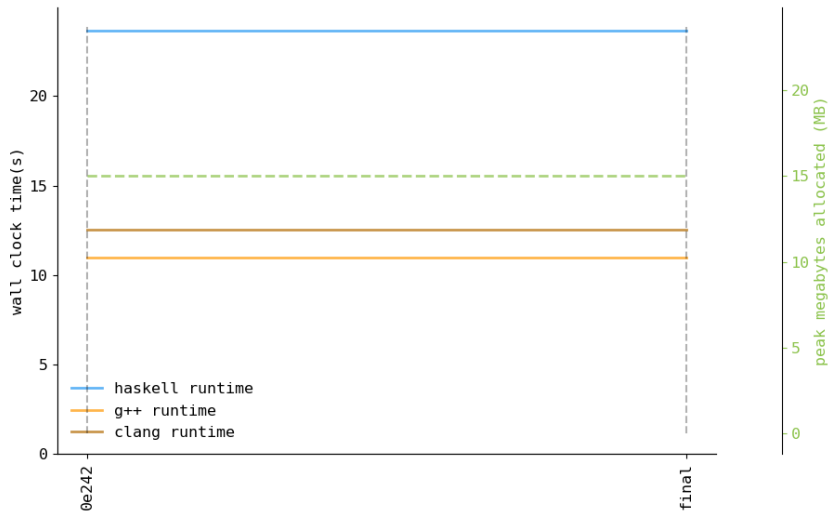
Initial Haskell Code: File I/O (1×)

```
withFile "image.ppm" WriteMode $ \hdl -> do
  hPrintf hdl "P3\n%d %d\n%d\n" w h (255::Int)
  flip mapM_ [0..w*h-1] \i -> do
    Vec r g b <- VM.unsafeRead c i
    hPrintf hdl "%d %d %d " (toInt r) (toInt g) (toInt b)
```


Initial Haskell Code: RNG (1×)

```
foreign import ccall unsafe "erand48"  
  erand48 :: Ptr Cushort -> IO Double
```

Performance: Initial Haskell Code



Restrict export list to main ($1\times \mapsto 1.13\times$)


```
-module Main where  
+module Main (main) where
```

Restrict export list to main ($1\times \mapsto 1.13\times$)

```
-module Main where  
+module Main (main) where
```


- ▶ Exported functions could be used by something unknown.
- ▶ Original versions must be available.

Performance: Restrict export list to main ($1\times \mapsto 1.13\times$)



Mark entries of Ray and Sphere as UNPACK and Strict ($1.13\times \mapsto 1.07\times$)



```
data Ray = Ray Vec Vec
```



- ▶ By default, all fields are *thunks* to rest of computation
- ▶ Pure, allow equational reasoning.

Mark entries of Ray and Sphere as UNPACK and Strict ($1.13\times \mapsto 1.07\times$)


```
data Ray = Ray !Vec !Vec
```



- ▶ When strict, elements are *pointers* to known structures
- ▶ pointers enable sharing!

Mark entries of Ray and Sphere as UNPACK and Strict ($1.13\times \mapsto 1.07\times$)

```
data Ray = Ray {-# UNPACK #-} !Vec {-# UNPACK #-} !Vec
```



- ▶ When unpacked, elements are *members* of the parent.
- ▶ Larger, but eliminate pointer chasing.

Mark entries of Ray and Sphere as UNPACK and Strict ($1.13\times \mapsto 1.07\times$)

```
data Vec = Vec {-# UNPACK #-} !Double
              {-# UNPACK #-} !Double
              {-# UNPACK #-} !Double


-data Ray = Ray Vec Vec -- origin, direction
+data Ray = Ray !Vec !Vec -- origin, direction

data Refl = DIFF | SPEC | REFR -- material types, used in radiance

-- radius, position, emission, color, reflection
-data Sphere = Sphere Double Vec Vec Vec !Refl
+data Sphere = Sphere {-# UNPACK #-} !Double
+                    {-# UNPACK #-} !Vec
+                    {-# UNPACK #-} !Vec
+                    {-# UNPACK #-} !Vec !Refl

struct Vec { double x, y, z; }
struct Ray { std::function<Vec()> v; std::function<Vec()> w; };
struct RayUnpack { double xv, yv, int zv;
                  double xw, yw, zw; };
```

Performance: Mark entries of Ray and Sphere as UNPACK and Strict
($1.13\times \mapsto 1.07\times$)



Use a pattern synonym to unpack Refl in Sphere ($1.07\times \mapsto 1.07\times$)

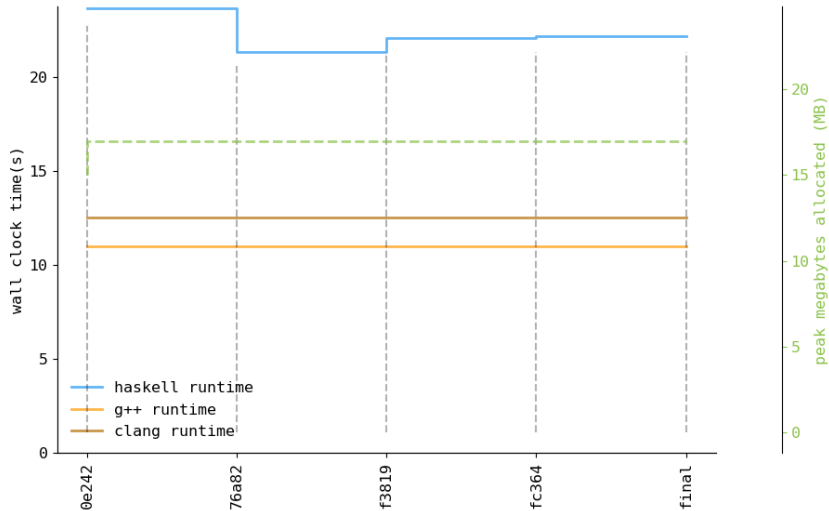
```
+{-# LANGUAGE PatternSynonyms #-}

-data Refl = DIFF | SPEC | REFR -- material types, used in radiance
+newtype Refl = Refl Int -- material types, used in radiance
+pattern DIFF,SPEC,REFR :: Refl
+pattern DIFF = Refl 0
+pattern SPEC = Refl 1
+pattern REFR = Refl 2
+{-# COMPLETE DIFF, SPEC, REFR #-}

-- radius, position, emission, color, reflection
data Sphere = Sphere {-# UNPACK #-} !Double
                    {-# UNPACK #-} !Vec
                    {-# UNPACK #-} !Vec
-                    {-# UNPACK #-} !Vec !Refl
+                    {-# UNPACK #-} !Vec {-# UNPACK #-} !Refl
```

Performance: Use a pattern synonym to unpack Refl in Sphere

($1.07\times \mapsto 1.07\times$)




Change from maximum on a list to max ($1.07\times \mapsto 1.08\times$)

```
-maxv (Vec a b c) = maximum [a,b,c]
+maxv (Vec a b c) = max a (max b c)

    let x = o `addv` (d `mulvs` t)
        n = norm $ x `subv` p
        nl = if n `dot` d < 0 then n else n `mulvs` (-1)
-      pr = maxv c
    depth' = depth + 1
    continue f = case refl of
      DIFF -> do
...
    if depth'>5
      then do
        er <- erand48 xi
+      let !pr = maxv c
```

Performance: Change from maximum on a list to max ($1.07\times \mapsto 1.08\times$)




Convert erand48 to pure Haskell ($1.08\times \mapsto 1.10\times$)

```
-foreign import ccall unsafe "erand48"
-  erand48 :: Ptr CUShort -> IO Double

+erand48 :: IORef Word64 -> IO Double
+erand48 !t = do -- | Some number crunching thing.
+  r <- readIORef t
+  let x' = 0x5deece66d * r + 0xb
+      d_word = 0x3ff0000000000000 .|. ((x' .&. 0xffffffffffff) `unsafeShiftL` 4)
+      d = castWord64ToDouble d_word - 1.0
+  writeIORef t x'
+  pure d
+...

-radiance :: Ray -> CInt -> Ptr CUShort -> IO Vec
+radiance :: Ray -> Int -> IORef Word64 -> IO Vec -- IORef with state
  radiance ray@(Ray o d) depth xi = case intersects ray of
    ...
    c <- VM.replicate (w * h) zeroV
  -  allocaArray 3 $ \xi -> -- Old RNG state
  -    flip mapM_ [0..h-1] $ \y -> do
+  xi <- newIORef 0 -- New RNG state
+  flip mapM_ [0..h-1] $ \y -> do
    writeXi xi y
```

Performance: Convert erand48 to pure Haskell ($1.08\times \mapsto 1.10\times$)




Remove mutability: Erand48 Monad ($1.10\times \mapsto 1.15\times$)

```
-erand48 :: IORef Word64 -> IO Double
-erand48 !t = do
-  r <- readIORef t
+data ET a = ET !Word64 !a deriving Functor
+newtype Erand48 a = Erand48 { runErand48' :: Word64 -> ET a } deriving Functor
+instance Applicative Erand48 where
+instance Monad Erand48 where
+runWithErand48 :: Int -> Erand48 a -> a
+erand48 :: Erand48 Double
...
-radiance :: Ray -> Int -> IORef Word64 -> IO Vec
-radiance ray@(Ray o d) depth xi = case intersects ray of
+radiance :: Ray -> Int -> Erand48 Vec
+radiance ray@(Ray o d) depth = case intersects ray of
...
-      r1 <- (2*pi*) <$> erand48 xi
-      r2 <- erand48 xi
+      r1 <- (2*pi*) <$> erand48
+      r2 <- erand48
...
-      then (* rp) <$> radiance reflRay depth' xi
-      else (* tp) <$> radiance (Ray x tdir) depth' xi
+      then (* rp) <$> radiance reflRay depth'
+      else (* tp) <$> radiance (Ray x tdir) depth'
```

Remove mutability: eliminate IORef and Data.Vector.Mutable (1.10× \mapsto 1.15×)

```
- c <- VM.replicate (w * h) 0
- xi <- newIORef 0
- flip mapM_ [0..h-1] $ \y -> do
-   writeXi xi y
-   for_ [0..w-1] \x -> do
-     let i = (h-y-1) * w + x
-     for_ [0..1] \sy -> do
-       for_ [0..1] \sx -> do
-         r <- newIORef 0
-         for_ [0..samps-1] \_s -> do
-           r1 <- (2*) <$> erand48 xi
+ img = (`concatMap` [(h-1),(h-2)..0]) $ \y -> runWithErand48 y do
+   for [0..w-1] \x -> do
+     (\pf -> foldlM pf 0 [(sy, sx) | sy <- [0,1], sx <- [0,1]]) \ci (sy, sx) -> do
+       Vec rr rg rb <- (\f -> foldlM f 0 [0..samps-1]) \ !r _s -> do
+         r1 <- (2*) <$> erand48
...
-     modifyIORef r (+ rad .* recip (fromIntegral samps))
-     ci <- VM.unsafeRead c i
-     Vec rr rg rb <- readIORef r
-     VM.unsafeWrite c i $ ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25
+     pure (r + rad .* recip (fromIntegral samps))
+     pure (ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25)
...
```

Performance: Remove mutability ($1.10\times \mapsto 1.15\times$)



Set everything in smallpt to be strict ($1.10\times \mapsto 1.15\times$)

```
smallpt :: Int -> Int -> Int -> IO ()
smallpt w h nsamps = do
  let !samps = nsamps `div` 4
      !org = Vec 50 52 295.6
      !dir = norm $ Vec 0 (-0.042612) (-1)
      !cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
      !cy = norm (cx `cross` dir) .* 0.5135
      !img = (`concatMap` [(h-1),(h-2)..0]) $ \y -> runWithErands48 y do
        for [0..w-1] \x -> do
          (\pf -> foldlM pf 0 [(sy, sx) | sy <- [0,1], sx <- [0,1]]) \ci (sy, sx) -> do
            !(Vec rr rg rb) <- (\f -> foldlM f 0 [0..samps-1]) \!r _s -> do
              !r1 <- (2*) <$> erand48
              let !dx = if r1<1 then sqrt r1-1 else 1-sqrt(2-r1)
              !r2 <- (2*) <$> erand48
              let !dy = if r2<1 then sqrt r2-1 else 1-sqrt(2-r2)
              !d = (cx .* (((sx + 0.5 + dx)/2 + fromIntegral x)/fromIntegral w - 0.5)) +
                  (cy .* (((sy + 0.5 + dy)/2 + fromIntegral y)/fromIntegral h - 0.5)) +
                  dir
              !rad <- radiance (Ray (org+d.*140) (norm d)) 0
            pure $! r + rad .* recip (fromIntegral samps)
          pure $! ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25
```

Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
let fooOpt = \y -> y
```

Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
let fooOpt = \y -> y
let foo' = let !x = error "ERR" in \y -> y
```

Why strictness may be bad

```
let foo = let x = error "ERR" in \y -> y
let fooOpt = \y -> y

let foo' = let !x = error "ERR" in \y -> y
let foo'Opt = \y -> y  -- ERROR! forcing foo'=foo'Opt should give "ERR"
```


Reduce to only useful strictnesses in smallpt ($1.10\times \mapsto 1.15\times$)

```
smallpt :: Int -> Int -> Int -> IO ()
smallpt w h nsamps = do
  let samps = nsamps `div` 4 -- NO LONGER STRICT
      org = Vec 50 52 295.6 -- NO LONGER STRICT
      dir = norm $ Vec 0 (-0.042612) (-1) -- NO LONGER STRICT
      cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
      cy = norm (cx `cross` dir) .* 0.5135
  img = (`concatMap` [(h-1),(h-2)..0]) $ \y -> runWithErand48 y do
    for [0..w-1] \x -> do
      (\pf -> foldlM pf 0 [(sy, sx) | sy <- [0,1], sx <- [0,1]]) \ci (sy, sx) -> do
        Vec rr rg rb <- (\f -> foldlM f 0 [0..samps-1]) \!r _s -> do
          r1 <- (2*) <$> erand48
          let !dx = if r1<1 then sqrt r1-1 else 1-sqrt(2-r1)
          r2 <- (2*) <$> erand48
          -- / STRICT
          let !dy = if r2<1 then sqrt r2-1 else 1-sqrt(2-r2)
              !d = (cx .* (((sx + 0.5 + dx)/2 + fromIntegral x)/fromIntegral w - 0.5)) +
                  (cy .* (((sy + 0.5 + dy)/2 + fromIntegral y)/fromIntegral h - 0.5)) +
                  dir
          rad <- radiance (Ray (org+d.*140) (norm d)) 0
          pure (r + rad .* recip (fromIntegral samps))
        pure (ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25)
```


Strategic application of strictness in entire project ($1.15\times \mapsto 1.23\times$)

```
...
- if det<0 then Nothing else f (b-sdet) (b+sdet)
- where op = p - o
-       eps = 1e-4
-       b = dot op d
-       det = b*b - dot op op + r*r
-       sdet = sqrt det
-       f a s = if a>eps then Just a else if s>eps then Just s else Nothing
+ if det<0
+ then Nothing
+ else
+   let !eps = 1e-4
+       !sdet = sqrt det
+       !a = b-sdet
+       !s = b+sdet
+   in if a>eps then Just a else if s>eps then Just s else Nothing
...
```

Performance: Strategic application of strictness in entire project ($1.15\times \mapsto 1.23\times$)




Remove Maybe from intersect(s) ($1.23\times \mapsto 1.40\times$)

```
| Old: Use Maybe Double to represent (was-hit?:bool, hit-distance: Double)
| New: use (1/0) to represent not (was-hit?)
-intersect :: Ray -> Sphere -> Maybe Double
+intersect :: Ray -> Sphere -> Double
intersect (Ray o d) (Sphere r p _e _c _refl) =
-  if det<0 then Nothing else f (b-sdet) (b+sdet)
+  if det<0 then (1/0.0) else f (b-sdet) (b+sdet)
  where op = p `subv` o
      ...
-      f a s = if a>eps then Just a else if s>eps then Just s else Nothing
+      f a s = if a>eps then a else if s>eps then s else (1/0.0)

-intersects :: Ray -> (Maybe Double, Sphere)
+intersects :: Ray -> (Double, Sphere)
intersects ray = (k, s)
-  where (k,s) = foldl' f (Nothing,undefined) spheres
-          f (k',sp) s' = case (k',intersect ray s') of
-              (Nothing,Just x) -> (Just x,s')
-              (Just y,Just x) | x < y -> (Just x,s')
-              _ -> (k',sp)
+  where (k,s) = foldl' f (1/0.0,undefined) spheres
+          f (k', sp) s' = let !x = intersect ray s' in if x < k' then (x, s') else (k', sp)

radiance :: Ray -> Int -> STRefU s Word64 -> ST s Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
-  (Nothing,_) -> return zerov
-  (Just t,Sphere _r p e c refl) -> do
+  (t,_) | t == (1/0.0) -> return zerov
+  (t,Sphere _r p e c refl) -> do
```

Performance: Remove Maybe from intersect(s) ($1.23\times \mapsto 1.40\times$)




Hand unroll the fold in intersects ($1.40\times \mapsto 1.43\times$)

```
intersects :: Ray -> (Double, Sphere)
-intersects ray = (k, s)
-  where (k,s) = foldl' f (1/0.0,undefined) spheres
+intersects ray =
+  f (... (f (f (intersect ray sphLeft, sphLeft) sphRight) ...))
+  where
+    f (k', sp) s' = let !x = intersect ray s' in if x < k' then (x, s') else (k', sp)

-spheres :: [Sphere]
-spheres = let s = Sphere ; z = zerov ; (.* ) = mulvs ; v = Vec in
-  [ s 1e5 (v (1e5+1) 40.8 81.6)    z (v 0.75 0.25 0.25) DIFF --Left
-    , s 1e5 (v (-1e5+99) 40.8 81.6) z (v 0.25 0.25 0.75) DIFF --Right
-  ...

+sphLeft, sphRight, ... :: Sphere
+sphLeft  = Sphere 1e5  (Vec (1e5+1) 40.8 81.6)    zerov (Vec 0.75 0.25 0.25) DIFF
+sphRight = Sphere 1e5  (Vec (-1e5+99) 40.8 81.6)  zerov (Vec 0.25 0.25 0.75) DIFF
+...
```

Performance: Hand unroll the fold in intersects ($1.40\times \mapsto 1.43\times$)



Custom datatype for intersects parameter passing ($1.43\times \mapsto 1.46\times$)


Old: Tuple with possibly-unevaluated Double and Sphere

New: Reference to a guaranteed-to-be-evaluated Double and Sphere

```
-intersects :: Ray -> (Double, Sphere)
+data T = T !Double !Sphere
+
+intersects :: Ray -> T
+  intersects ray =
-    f ( ... f (intersect ray sphLeft, sphLeft) sphRight) ... sphLite
+    f ( ... f (T (intersect ray sphLeft) sphLeft) sphRight) ... sphLite
+  where
-    f (k', sp) s' =
-      let !x = intersect ray s' in if x < k' then (x, s') else (k', sp)
+    f !(T k' sp) !s' =
+      let !x = intersect ray s' in if x < k' then T x s' else T k' sp

radiance :: Ray -> Int -> Erand48 Vec
radiance ray@(Ray o d) depth = case intersects ray of
-  (!t,_) | t == 1/0.0 -> return 0
-  (!t,!Sphere _r p e c refl) -> do
+  (T t _) | t == 1/0.0 -> return 0
+  (T t (Sphere _r p e c refl)) -> do
    let !x = o + d .* t
        !n = norm $ x - p
        !nl = if dot n d < 0 then n else negate n
```

Performance: Custom datatype for intersects parameter passing ($1.43\times \mapsto 1.46\times$)




Optimize file writing: ($1.46\times \mapsto 1.46\times$)

```
build-depends:
    base >= 4.12 && < 4.15
+   , bytestring ^>= 0.11

-toInt :: Double -> Int
-toInt x = floor $ clamp x ** recip 2.2 * 255 + 0.5
+toInt :: Double -> BB.Builder -- 0(1) concatenation
+toInt x = BB.intDec (floor (clamp x ** recip 2.2 * 255 + 0.5)) <> BB.char8 ' '
...
withFile "image.ppm" WriteMode $ \hdl -> do
-   hPrintf hdl "P3\n%d %d\n%d\n" w h (255::Int)
-   for_ img \(Vec r g b) -> do
-       hPrintf hdl "%d %d %d " (toInt r) (toInt g) (toInt b)
+   BB.hPutBuilder hdl $
+       BB.string8 "P3\n" <> -- efficient builders for ASCII
+       BB.intDec w <> BB.char8 ' ' <> BB.intDec h <> BB.char8 '\n' <>
+       BB.intDec 255 <> BB.char8 '\n' <>
+       (mconcat $ fmap \(Vec r g b) -> toInt r <> toInt g <> toInt b) img)
```


Performance: Optimize file writing ($1.46\times \mapsto 1.46\times$)



Use LLVM backend ($1.46\times \mapsto 2.04\times$)

```
+package smallpt-opt  
+ ghc-options: -fllvm
```

The view from the mountaintop ($1.46\times \mapsto 2.04\times$)



Avoid CPU ieee754 slow paths ($2.04\times \mapsto 2.12\times$)


```
intersect :: Ray -> Sphere -> Double
intersect (Ray o d) (Sphere r p _e _c _refl) =
    if det<0
-   then 1/0.0
+   then 1e20
    else
        ...
-   in if a>eps then a else if s>eps then s else 1/0.0
+   in if a>eps then a else if s>eps then s else 1e20
...

radiance :: Ray -> Int -> Erand48 Vec
radiance ray@(Ray o d) depth = case intersects ray of
-   (T t _) | t == 1/0.0 -> return 0
+   (T 1e20 _) -> return 0
...
```

Fix differences with C++ version ($2.12\times \mapsto 2.32\times$)

```
-         if depth>2  
+         if depth'>2 -- depth' = depth + 1  
...
```

A second view from the mountaintop



Takeaways

- ▶ The unrolling in 'intersects' is ugly.
- ▶ (We feel) the maintainability of this code hasn't been significantly harmed.
- ▶ We're faster than `clang++` and `g++`
- ▶ Haven't exhausted the optimization opportunities.
- ▶ GHC could learn to do several of these optimizations for us.
- ▶ Others are just good Haskell style.
- ▶ Clean Haskell is often performant Haskell.
- ▶ Repository stepping through each optimization is available at github.com/bollu/smallpt-opt
- ▶ Slides at github.com/bollu/slides-haskell-exchange-2020-smallpt

Raw data

- ▶ All test were on an otherwise idle Equinix Metal c3.small.x86 (Intel Xeon E-2278G with 32GiB RAM, Ubuntu 20.04). Averages over ten runs were reported. Performance numbers are stable.

0e242	23.6586	23.658	23.7251	23.6954	23.676	23.656	23.673	23.7139	23.6598	23.7641
76a82	21.3573	21.384	21.3516	21.3304	21.3658	21.3776	21.3564	21.3843	21.3401	21.3683
f3819	22.1036	22.0754	22.1034	22.0864	22.0692	22.1004	22.0584	22.1181	22.0728	22.0972
fc364	22.1211	22.114	22.1205	22.1101	23.0621	22.1101	22.1163	22.133	22.1491	22.1464
66c39	21.8626	21.8684	21.8977	21.9043	21.893	21.8483	21.8335	21.8869	21.8848	21.8335
5704a	21.4682	21.4692	21.4411	21.473	21.4783	21.4613	21.4818	21.4507	21.4388	21.4933
eb9a3	20.6134	20.6014	20.6527	20.596	20.6034	20.6174	20.5965	20.594	20.5967	20.5892
5caf8	20.5209	20.535	20.5312	20.5289	20.5338	20.5717	20.5387	20.5386	20.5262	20.5488
361eb	20.5551	20.5485	20.5602	20.552	20.5668	20.555	20.5573	20.5564	20.5599	20.5823
7911c	19.2532	19.2664	19.2955	19.2565	19.2517	19.2722	19.3284	19.2611	19.2623	19.2596
149bc	16.8551	16.8551	16.8788	16.9067	16.8683	16.8685	16.8651	16.9186	16.8589	16.8604
80931	16.5752	16.5739	16.5831	16.5918	16.5678	16.5785	16.6128	16.5682	16.5816	16.577
07d68	16.1819	16.1672	16.1829	16.2267	16.1716	16.1854	16.1806	16.1949	16.1917	16.1784
4e141	16.2206	16.1816	16.2002	16.1799	16.1813	16.1781	16.1929	16.243	16.1705	16.1877
16c56	11.6334	11.6166	11.6837	11.6504	11.6227	11.6135	11.5949	11.5966	11.6013	11.64
d6347	11.1632	11.218	11.1741	11.1802	11.1849	11.1755	11.1729	11.2155	11.1718	11.2089
9ee06	10.2131	10.2154	10.1994	10.2105	10.2028	10.2344	10.2008	10.2497	10.2226	10.3042
gcc	10.97	10.97	10.97	10.99	10.98	10.97	10.97	10.98	10.98	10.98
clang	12.53	12.51	12.5	12.53	12.51	12.5	12.52	12.48	12.53	12.52

- ▶ We're likely faster than C++ because we can see `erand48` that the C++ compiler cannot.
- ▶ We're also likely faster than C++ because we stream the pixel outputs that's written to the file.
- ▶ C++ keeps all of the pixel data resident in-memory.