

Optimizing smallpt

Davean Scies, Siddharth Bhat

Haskell Exchange

November 4th, 2020


What is `smallpt` anyway?

Optimizing smallpt

2020-11-04

└ What is smallpt anyway?


What is smallpt anyway?



Optimizing smallpt


2020-11-04

What is smallpt anyway?




- ▶ 100 LoC C demo of a raytracer

└─What is smallpt anyway?



- ▶ 100 LoC C demo of a raytracer

What is smallpt anyway?




- ▶ 100 LoC C demo of a raytracer
- ▶ Perfect for an optimization case study

2020-11-04

Optimizing smallpt

What is smallpt anyway?



- ▶ 100 LoC C demo of a raytracer
- ▶ Perfect for an optimization case study

Note for first slide: what is smallpt anyway?

What is smallpt anyway?

```
struct Vec {
    double x, y, z; // position, also color (r,g,b)
    ... methods...
};

struct Ray { Vec o, d; Ray(Vec o_, Vec d_) : o(o_), d(d_) {} };

enum Refl_t { DIFF, SPEC, REFR }; // material types, used in radiance()

struct Sphere {
    double rad; // radius
    Vec p, e, c; // position, emission, color
    Refl_t refl; // reflection type (DIFFuse, SPECular, REFRactive)
    ... methods ...
    double intersect(const Ray &r) const // returns distance, 0 if nohit
};

Sphere spheres[] = { //Scene: radius, position, emission, color, material
    Sphere(1e5, Vec( 1e5+1,40.8,81.6), Vec(),Vec(.75,.25,.25),DIFF), //Left
    ... initialization ...
};
```

Optimizing smallpt

2020-11-04

What is smallpt anyway?

What is smallpt anyway?

```
struct Vec {
    double x, y, z; // position, also color (r,g,b)
    ... methods...
};

struct Ray { Vec o, d; Ray(Vec o_, Vec d_) : o(o_), d(d_) {} };

enum Refl_t { DIFF, SPEC, REFR }; // material types, used in radiance()

struct Sphere {
    double rad; // radius
    Vec p, e, c; // position, emission, color
    Refl_t refl; // reflection type (DIFFuse, SPECular, REFRactive)
    ... methods ...
    double intersect(const Ray &r) const // returns distance, 0 if nohit
};

Sphere spheres[] = { //Scene: radius, position, emission, color, material
    Sphere(1e5, Vec( 1e5+1,40.8,81.6), Vec(),Vec(.75,.25,.25),DIFF), //Left
    ... initialization ...
};
```

What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
```

```
}
```

Optimizing smallpt

2020-11-04

What is smallpt anyway?

Say that the core function is radiance



What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
```

radiance

radiance

radiance

```
radiance radiance  
radiance radiance  
}
```

Optimizing smallpt

2020-11-04

What is smallpt anyway?



What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
```

```
    if ( /* */ ) if ( /* */ ) else  
    if ( /* */ ){
```

```
        radiance  
    } else if ( /* */ )  
        radiance
```

```
    if ( /* */ )  
        radiance
```

```
    radiance  
    radiance  
}
```

Optimizing smallpt

2020-11-04

What is smallpt anyway?

```
What is smallpt anyway?  
  
Vec radiance(const Ray &r, int depth, unsigned short *Xi){  
  
    if ( /* */ ) if ( /* */ ) else  
    if ( /* */ ){  
  
        radiance  
    } else if ( /* */ )  
        radiance  
  
    if ( /* */ )  
        radiance  
  
    radiance  
    radiance  
}
```

...with control flow

What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
```

```
Vec x=r.o+r.d*t, n=(x-obj.p).norm(), nl=n.dot(r.d)<0?n:n*-1, f=obj.c;
```

```
if ( ) if ( ) else  
if ( ){
```

```
radiance  
} else if ( )  
radiance
```

```
if ((cos2t=1-nnt*nnt*(1-ddn*ddn))<0)  
radiance
```

```
radiance radiance  
radiance radiance  
}
```

Optimizing smallpt

2020-11-04

What is smallpt anyway?

and lots of arithmetic

```
What is smallpt anyway?  
  
Vec radiance(const Ray &r, int depth, unsigned short *Xi){  
  
Vec w=obj.a*0.5, w1=obj.g*0.5, w2=obj.b*0.5, w3=obj.c*0.5, w4=obj.d*0.5, w5=obj.e*0.5, w6=obj.f*0.5, w7=obj.g*0.5, w8=obj.h*0.5, w9=obj.i*0.5, w10=obj.j*0.5, w11=obj.k*0.5, w12=obj.l*0.5, w13=obj.m*0.5, w14=obj.n*0.5, w15=obj.o*0.5, w16=obj.p*0.5, w17=obj.q*0.5, w18=obj.r*0.5, w19=obj.s*0.5, w20=obj.t*0.5, w21=obj.u*0.5, w22=obj.v*0.5, w23=obj.w*0.5, w24=obj.x*0.5, w25=obj.y*0.5, w26=obj.z*0.5, w27=obj.aa*0.5, w28=obj.bb*0.5, w29=obj.cc*0.5, w30=obj.dd*0.5, w31=obj.ee*0.5, w32=obj.ff*0.5, w33=obj.gg*0.5, w34=obj.hh*0.5, w35=obj.ii*0.5, w36=obj.jj*0.5, w37=obj.kk*0.5, w38=obj.ll*0.5, w39=obj.mm*0.5, w40=obj.nn*0.5, w41=obj.oo*0.5, w42=obj.pp*0.5, w43=obj.qq*0.5, w44=obj.rr*0.5, w45=obj.ss*0.5, w46=obj.tt*0.5, w47=obj.uu*0.5, w48=obj.vv*0.5, w49=obj.ww*0.5, w50=obj.xx*0.5, w51=obj.yy*0.5, w52=obj.zz*0.5, w53=obj.aa*0.5, w54=obj.bb*0.5, w55=obj.cc*0.5, w56=obj.dd*0.5, w57=obj.ee*0.5, w58=obj.ff*0.5, w59=obj.gg*0.5, w60=obj.hh*0.5, w61=obj.ii*0.5, w62=obj.jj*0.5, w63=obj.kk*0.5, w64=obj.ll*0.5, w65=obj.mm*0.5, w66=obj.nn*0.5, w67=obj.oo*0.5, w68=obj.pp*0.5, w69=obj.qq*0.5, w70=obj.rr*0.5, w71=obj.ss*0.5, w72=obj.tt*0.5, w73=obj.uu*0.5, w74=obj.vv*0.5, w75=obj.ww*0.5, w76=obj.xx*0.5, w77=obj.yy*0.5, w78=obj.zz*0.5, w79=obj.aa*0.5, w80=obj.bb*0.5, w81=obj.cc*0.5, w82=obj.dd*0.5, w83=obj.ee*0.5, w84=obj.ff*0.5, w85=obj.gg*0.5, w86=obj.hh*0.5, w87=obj.ii*0.5, w88=obj.jj*0.5, w89=obj.kk*0.5, w90=obj.ll*0.5, w91=obj.mm*0.5, w92=obj.nn*0.5, w93=obj.oo*0.5, w94=obj.pp*0.5, w95=obj.qq*0.5, w96=obj.rr*0.5, w97=obj.ss*0.5, w98=obj.tt*0.5, w99=obj.uu*0.5, w100=obj.vv*0.5, w101=obj.ww*0.5, w102=obj.xx*0.5, w103=obj.yy*0.5, w104=obj.zz*0.5, w105=obj.aa*0.5, w106=obj.bb*0.5, w107=obj.cc*0.5, w108=obj.dd*0.5, w109=obj.ee*0.5, w110=obj.ff*0.5, w111=obj.gg*0.5, w112=obj.hh*0.5, w113=obj.ii*0.5, w114=obj.jj*0.5, w115=obj.kk*0.5, w116=obj.ll*0.5, w117=obj.mm*0.5, w118=obj.nn*0.5, w119=obj.oo*0.5, w120=obj.pp*0.5, w121=obj.qq*0.5, w122=obj.rr*0.5, w123=obj.ss*0.5, w124=obj.tt*0.5, w125=obj.uu*0.5, w126=obj.vv*0.5, w127=obj.ww*0.5, w128=obj.xx*0.5, w129=obj.yy*0.5, w130=obj.zz*0.5, w131=obj.aa*0.5, w132=obj.bb*0.5, w133=obj.cc*0.5, w134=obj.dd*0.5, w135=obj.ee*0.5, w136=obj.ff*0.5, w137=obj.gg*0.5, w138=obj.hh*0.5, w139=obj.ii*0.5, w140=obj.jj*0.5, w141=obj.kk*0.5, w142=obj.ll*0.5, w143=obj.mm*0.5, w144=obj.nn*0.5, w145=obj.oo*0.5, w146=obj.pp*0.5, w147=obj.qq*0.5, w148=obj.rr*0.5, w149=obj.ss*0.5, w150=obj.tt*0.5, w151=obj.uu*0.5, w152=obj.vv*0.5, w153=obj.ww*0.5, w154=obj.xx*0.5, w155=obj.yy*0.5, w156=obj.zz*0.5, w157=obj.aa*0.5, w158=obj.bb*0.5, w159=obj.cc*0.5, w160=obj.dd*0.5, w161=obj.ee*0.5, w162=obj.ff*0.5, w163=obj.gg*0.5, w164=obj.hh*0.5, w165=obj.ii*0.5, w166=obj.jj*0.5, w167=obj.kk*0.5, w168=obj.ll*0.5, w169=obj.mm*0.5, w170=obj.nn*0.5, w171=obj.oo*0.5, w172=obj.pp*0.5, w173=obj.qq*0.5, w174=obj.rr*0.5, w175=obj.ss*0.5, w176=obj.tt*0.5, w177=obj.uu*0.5, w178=obj.vv*0.5, w179=obj.ww*0.5, w180=obj.xx*0.5, w181=obj.yy*0.5, w182=obj.zz*0.5, w183=obj.aa*0.5, w184=obj.bb*0.5, w185=obj.cc*0.5, w186=obj.dd*0.5, w187=obj.ee*0.5, w188=obj.ff*0.5, w189=obj.gg*0.5, w190=obj.hh*0.5, w191=obj.ii*0.5, w192=obj.jj*0.5, w193=obj.kk*0.5, w194=obj.ll*0.5, w195=obj.mm*0.5, w196=obj.nn*0.5, w197=obj.oo*0.5, w198=obj.pp*0.5, w199=obj.qq*0.5, w200=obj.rr*0.5, w201=obj.ss*0.5, w202=obj.tt*0.5, w203=obj.uu*0.5, w204=obj.vv*0.5, w205=obj.ww*0.5, w206=obj.xx*0.5, w207=obj.yy*0.5, w208=obj.zz*0.5, w209=obj.aa*0.5, w210=obj.bb*0.5, w211=obj.cc*0.5, w212=obj.dd*0.5, w213=obj.ee*0.5, w214=obj.ff*0.5, w215=obj.gg*0.5, w216=obj.hh*0.5, w217=obj.ii*0.5, w218=obj.jj*0.5, w219=obj.kk*0.5, w220=obj.ll*0.5, w221=obj.mm*0.5, w222=obj.nn*0.5, w223=obj.oo*0.5, w224=obj.pp*0.5, w225=obj.qq*0.5, w226=obj.rr*0.5, w227=obj.ss*0.5, w228=obj.tt*0.5, w229=obj.uu*0.5, w230=obj.vv*0.5, w231=obj.ww*0.5, w232=obj.xx*0.5, w233=obj.yy*0.5, w234=obj.zz*0.5, w235=obj.aa*0.5, w236=obj.bb*0.5, w237=obj.cc*0.5, w238=obj.dd*0.5, w239=obj.ee*0.5, w240=obj.ff*0.5, w241=obj.gg*0.5, w242=obj.hh*0.5, w243=obj.ii*0.5, w244=obj.jj*0.5, w245=obj.kk*0.5, w246=obj.ll*0.5, w247=obj.mm*0.5, w248=obj.nn*0.5, w249=obj.oo*0.5, w250=obj.pp*0.5, w251=obj.qq*0.5, w252=obj.rr*0.5, w253=obj.ss*0.5, w254=obj.tt*0.5, w255=obj.uu*0.5, w256=obj.vv*0.5, w257=obj.ww*0.5, w258=obj.xx*0.5, w259=obj.yy*0.5, w260=obj.zz*0.5, w261=obj.aa*0.5, w262=obj.bb*0.5, w263=obj.cc*0.5, w264=obj.dd*0.5, w265=obj.ee*0.5, w266=obj.ff*0.5, w267=obj.gg*0.5, w268=obj.hh*0.5, w269=obj.ii*0.5, w270=obj.jj*0.5, w271=obj.kk*0.5, w272=obj.ll*0.5, w273=obj.mm*0.5, w274=obj.nn*0.5, w275=obj.oo*0.5, w276=obj.pp*0.5, w277=obj.qq*0.5, w278=obj.rr*0.5, w279=obj.ss*0.5, w280=obj.tt*0.5, w281=obj.uu*0.5, w282=obj.vv*0.5, w283=obj.ww*0.5, w284=obj.xx*0.5, w285=obj.yy*0.5, w286=obj.zz*0.5, w287=obj.aa*0.5, w288=obj.bb*0.5, w289=obj.cc*0.5, w290=obj.dd*0.5, w291=obj.ee*0.5, w292=obj.ff*0.5, w293=obj.gg*0.5, w294=obj.hh*0.5, w295=obj.ii*0.5, w296=obj.jj*0.5, w297=obj.kk*0.5, w298=obj.ll*0.5, w299=obj.mm*0.5, w300=obj.nn*0.5, w301=obj.oo*0.5, w302=obj.pp*0.5, w303=obj.qq*0.5, w304=obj.rr*0.5, w305=obj.ss*0.5, w306=obj.tt*0.5, w307=obj.uu*0.5, w308=obj.vv*0.5, w309=obj.ww*0.5, w310=obj.xx*0.5, w311=obj.yy*0.5, w312=obj.zz*0.5, w313=obj.aa*0.5, w314=obj.bb*0.5, w315=obj.cc*0.5, w316=obj.dd*0.5, w317=obj.ee*0.5, w318=obj.ff*0.5, w319=obj.gg*0.5, w320=obj.hh*0.5, w321=obj.ii*0.5, w322=obj.jj*0.5, w323=obj.kk*0.5, w324=obj.ll*0.5, w325=obj.mm*0.5, w326=obj.nn*0.5, w327=obj.oo*0.5, w328=obj.pp*0.5, w329=obj.qq*0.5, w330=obj.rr*0.5, w331=obj.ss*0.5, w332=obj.tt*0.5, w333=obj.uu*0.5, w334=obj.vv*0.5, w335=obj.ww*0.5, w336=obj.xx*0.5, w337=obj.yy*0.5, w338=obj.zz*0.5, w339=obj.aa*0.5, w340=obj.bb*0.5, w341=obj.cc*0.5, w342=obj.dd*0.5, w343=obj.ee*0.5, w344=obj.ff*0.5, w345=obj.gg*0.5, w346=obj.hh*0.5, w347=obj.ii*0.5, w348=obj.jj*0.5, w349=obj.kk*0.5, w350=obj.ll*0.5, w351=obj.mm*0.5, w352=obj.nn*0.5, w353=obj.oo*0.5, w354=obj.pp*0.5, w355=obj.qq*0.5, w356=obj.rr*0.5, w357=obj.ss*0.5, w358=obj.tt*0.5, w359=obj.uu*0.5, w360=obj.vv*0.5, w361=obj.ww*0.5, w362=obj.xx*0.5, w363=obj.yy*0.5, w364=obj.zz*0.5, w365=obj.aa*0.5, w366=obj.bb*0.5, w367=obj.cc*0.5, w368=obj.dd*0.5, w369=obj.ee*0.5, w370=obj.ff*0.5, w371=obj.gg*0.5, w372=obj.hh*0.5, w373=obj.ii*0.5, w374=obj.jj*0.5, w375=obj.kk*0.5, w376=obj.ll*0.5, w377=obj.mm*0.5, w378=obj.nn*0.5, w379=obj.oo*0.5, w380=obj.pp*0.5, w381=obj.qq*0.5, w382=obj.rr*0.5, w383=obj.ss*0.5, w384=obj.tt*0.5, w385=obj.uu*0.5, w386=obj.vv*0.5, w387=obj.ww*0.5, w388=obj.xx*0.5, w389=obj.yy*0.5, w390=obj.zz*0.5, w391=obj.aa*0.5, w392=obj.bb*0.5, w393=obj.cc*0.5, w394=obj.dd*0.5, w395=obj.ee*0.5, w396=obj.ff*0.5, w397=obj.gg*0.5, w398=obj.hh*0.5, w399=obj.ii*0.5, w400=obj.jj*0.5, w401=obj.kk*0.5, w402=obj.ll*0.5, w403=obj.mm*0.5, w404=obj.nn*0.5, w405=obj.oo*0.5, w406=obj.pp*0.5, w407=obj.qq*0.5, w408=obj.rr*0.5, w409=obj.ss*0.5, w410=obj.tt*0.5, w411=obj.uu*0.5, w412=obj.vv*0.5, w413=obj.ww*0.5, w414=obj.xx*0.5, w415=obj.yy*0.5, w416=obj.zz*0.5, w417=obj.aa*0.5, w418=obj.bb*0.5, w419=obj.cc*0.5, w420=obj.dd*0.5, w421=obj.ee*0.5, w422=obj.ff*0.5, w423=obj.gg*0.5, w424=obj.hh*0.5, w425=obj.ii*0.5, w426=obj.jj*0.5, w427=obj.kk*0.5, w428=obj.ll*0.5, w429=obj.mm*0.5, w430=obj.nn*0.5, w431=obj.oo*0.5, w432=obj.pp*0.5, w433=obj.qq*0.5, w434=obj.rr*0.5, w435=obj.ss*0.5, w436=obj.tt*0.5, w437=obj.uu*0.5, w438=obj.vv*0.5, w439=obj.ww*0.5, w440=obj.xx*0.5, w441=obj.yy*0.5, w442=obj.zz*0.5, w443=obj.aa*0.5, w444=obj.bb*0.5, w445=obj.cc*0.5, w446=obj.dd*0.5, w447=obj.ee*0.5, w448=obj.ff*0.5, w449=obj.gg*0.5, w450=obj.hh*0.5, w451=obj.ii*0.5, w452=obj.jj*0.5, w453=obj.kk*0.5, w454=obj.ll*0.5, w455=obj.mm*0.5, w456=obj.nn*0.5, w457=obj.oo*0.5, w458=obj.pp*0.5, w459=obj.qq*0.5, w460=obj.rr*0.5, w461=obj.ss*0.5, w462=obj.tt*0.5, w463=obj.uu*0.5, w464=obj.vv*0.5, w465=obj.ww*0.5, w466=obj.xx*0.5, w467=obj.yy*0.5, w468=obj.zz*0.5, w469=obj.aa*0.5, w470=obj.bb*0.5, w471=obj.cc*0.5, w472=obj.dd*0.5, w473=obj.ee*0.5, w474=obj.ff*0.5, w475=obj.gg*0.5, w476=obj.hh*0.5, w477=obj.ii*0.5, w478=obj.jj*0.5, w479=obj.kk*0.5, w480=obj.ll*0.5, w481=obj.mm*0.5, w482=obj.nn*0.5, w483=obj.oo*0.5, w484=obj.pp*0.5, w485=obj.qq*0.5, w486=obj.rr*0.5, w487=obj.ss*0.5, w488=obj.tt*0.5, w489=obj.uu*0.5, w490=obj.vv*0.5, w491=obj.ww*0.5, w492=obj.xx*0.5, w493=obj.yy*0.5, w494=obj.zz*0.5, w495=obj.aa*0.5, w496=obj.bb*0.5, w497=obj.cc*0.5, w498=obj.dd*0.5, w499=obj.ee*0.5, w500=obj.ff*0.5, w501=obj.gg*0.5, w502=obj.hh*0.5, w503=obj.ii*0.5, w504=obj.jj*0.5, w505=obj.kk*0.5, w506=obj.ll*0.5, w507=obj.mm*0.5, w508=obj.nn*0.5, w509=obj.oo*0.5, w510=obj.pp*0.5, w511=obj.qq*0.5, w512=obj.rr*0.5, w513=obj.ss*0.5, w514=obj.tt*0.5, w515=obj.uu*0.5, w516=obj.vv*0.5, w517=obj.ww*0.5, w518=obj.xx*0.5, w519=obj.yy*0.5, w520=obj.zz*0.5, w521=obj.aa*0.5, w522=obj.bb*0.5, w523=obj.cc*0.5, w524=obj.dd*0.5, w525=obj.ee*0.5, w526=obj.ff*0.5, w527=obj.gg*0.5, w528=obj.hh*0.5, w529=obj.ii*0.5, w530=obj.jj*0.5, w531=obj.kk*0.5, w532=obj.ll*0.5, w533=obj.mm*0.5, w534=obj.nn*0.5, w535=obj.oo*0.5, w536=obj.pp*0.5, w537=obj.qq*0.5, w538=obj.rr*0.5, w539=obj.ss*0.5, w540=obj.tt*0.5, w541=obj.uu*0.5, w542=obj.vv*0.5, w543=obj.ww*0.5, w544=obj.xx*0.5, w545=obj.yy*0.5, w546=obj.zz*0.5, w547=obj.aa*0.5, w548=obj.bb*0.5, w549=obj.cc*0.5, w550=obj.dd*0.5, w551=obj.ee*0.5, w552=obj.ff*0.5, w553=obj.gg*0.5, w554=obj.hh*0.5, w555=obj.ii*0.5, w556=obj.jj*0.5, w557=obj.kk*0.5, w558=obj.ll*0.5, w559=obj.mm*0.5, w560=obj.nn*0.5, w561=obj.oo*0.5, w562=obj.pp*0.5, w563=obj.qq*0.5, w564=obj.rr*0.5, w565=obj.ss*0.5, w566=obj.tt*0.5, w567=obj.uu*0.5, w568=obj.vv*0.5, w569=obj.ww*0.5, w570=obj.xx*0.5, w571=obj.yy*0.5, w572=obj.zz*0.5, w573=obj.aa*0.5, w574=obj.bb*0.5, w575=obj.cc*0.5, w576=obj.dd*0.5, w577=obj.ee*0.5, w578=obj.ff*0.5, w579=obj.gg*0.5, w580=obj.hh*0.5, w581=obj.ii*0.5, w582=obj.jj*0.5, w583=obj.kk*0.5, w584=obj.ll*0.5, w585=obj.mm*0.5, w586=obj.nn*0.5, w587=obj.oo*0.5, w588=obj.pp*0.5, w589=obj.qq*0.5, w590=obj.rr*0.5, w591=obj.ss*0.5, w592=obj.tt*0.5, w593=obj.uu*0.5, w594=obj.vv*0.5, w595=obj.ww*0.5, w596=obj.xx*0.5, w597=obj.yy*0.5, w598=obj.zz*0.5, w599=obj.aa*0.5, w600=obj.bb*0.5, w601=obj.cc*0.5, w602=obj.dd*0.5, w603=obj.ee*0.5, w604=obj.ff*0.5, w605=obj.gg*0.5, w606=obj.hh*0.5, w607=obj.ii*0.5, w608=obj.jj*0.5, w609=obj.kk*0.5, w610=obj.ll*0.5, w611=obj.mm*0.5, w612=obj.nn*0.5, w613=obj.oo*0.5, w614=obj.pp*0.5, w615=obj.qq*0.5, w616=obj.rr*0.5, w617=obj.ss*0.5, w618=obj.tt*0.5, w619=obj.uu*0.5, w620=obj.vv*0.5, w621=obj.ww*0.5, w622=obj.xx*0.5, w623=obj.yy*0.5, w624=obj.zz*0.5, w625=obj.aa*0.5, w626=obj.bb*0.5, w627=obj.cc*0.5, w628=obj.dd*0.5, w629=obj.ee*0.5, w630=obj.ff*0.5, w631=obj.gg*0.5, w632=obj.hh*0.5, w633=obj.ii*0.5, w634=obj.jj*0.5, w635=obj.kk*0.5, w636=obj.ll*0.5, w637=obj.mm*0.5, w638=obj.nn*0.5, w639=obj.oo*0.5, w640=obj.pp*0.5, w641=obj.qq*0.5, w642=obj.rr*0.5, w643=obj.ss*0.5, w644=obj.tt*0.5, w645=obj.uu*0.5, w646=obj.vv*0.5, w647=obj.ww*0.5, w648=obj.xx*0.5, w649=obj.yy*0.5, w650=obj.zz*0.5, w651=obj.aa*0.5, w652=obj.bb*0.5, w653=obj.cc*0.5, w654=obj.dd*0.5, w655=obj.ee*0.5, w656=obj.ff*0.5, w657=obj.gg*0.5, w658=obj.hh*0.5, w659=obj.ii*0.5, w660=obj.jj*0.5, w661=obj.kk*0.5, w662=obj.ll*0.5, w663=obj.mm*0.5, w664=obj.nn*0.5, w665=obj.oo*0.5, w666=obj.pp*0.5, w667=obj.qq*0.5, w668=obj.rr*0.5, w669=obj.ss*0.5, w670=obj.tt*0.5, w671=obj.uu*0.5, w672=obj.vv*0.5, w673=obj.ww*0.5, w674=obj.xx*0.5, w675=obj.yy*0.5, w676=obj.zz*0.5, w677=obj.aa*0.5, w678=obj.bb*0.5, w679=obj.cc*0.5, w680=obj.dd*0.5, w681=obj.ee*0.5, w682=obj.ff*0.5, w683=obj.gg*0.5, w684=obj.hh*0.5, w685=obj.ii*0.5, w686=obj.jj*0.5, w687=obj.kk*0.5, w688=obj.ll*0.5, w689=obj.mm*0.5, w690=obj.nn*0.5, w691=obj.oo*0.5, w692=obj.pp*0.5, w693=obj.qq*0.5, w694=obj.rr*0.5, w695=obj.ss*0.5, w696=obj.tt*0.5, w697=obj.uu*0.5, w698=obj.vv*0.5, w699=obj.ww*0.5, w700=obj.xx*0.5, w701=obj.yy*0.5, w702=obj.zz*0.5, w703=obj.aa*0.5, w704=obj.bb*0.5, w705=obj.cc*0.5, w706=obj.dd*0.5, w707=obj.ee*0.5, w708=obj.ff*0.5, w709=obj.gg*0.5, w710=obj.hh*0.5, w711=obj.ii*0.5, w712=obj.jj*0.5, w713=obj.kk*0.5, w714=obj.ll*0.5, w715=obj.mm*0.5, w716=obj.nn*0.5, w717=obj.oo*0.5, w718=obj.pp*0.5, w719=obj.qq*0.5, w720=obj.rr*0.5, w721=obj.ss*0.5, w722=obj.tt*0.5, w723=obj.uu*0.5, w724=obj.vv*0.5, w725=obj.ww*0.5, w726=obj.xx*0.5, w727=obj.yy*0.5, w728=obj.zz*0.5, w729=obj.aa*0.5, w730=obj.bb*0.5, w731=obj.cc*0.5, w732=obj.dd*0.5, w733=obj.ee*0.5, w734=obj.ff*0.5, w735=obj.gg*0.5, w736=obj.hh*0.5, w737=obj.ii*0.5, w738=obj.jj*0.5, w739=obj.kk*0.5, w740=obj.ll*0.5, w741=obj.mm*0.5, w742=obj.nn*0.5, w743=obj.oo*0.5, w744=obj.pp*0.5, w745=obj.qq*0.5, w746=obj.rr*0.5, w747=obj.ss*0.5, w748=obj.tt*0.5, w749=obj.uu*0.5, w750=obj.vv*0.5, w751=obj.ww*0.5, w752=obj.xx*0.5, w753=obj.yy*0.5, w754=obj.zz*0.5, w755=obj.aa*0.5, w756=obj.bb*0.5, w757=obj.cc*0.5, w758=obj.dd*0.5, w759=obj.ee*0.5, w760=obj.ff*0.5, w761=obj.gg*0.5, w762=obj.hh*0.5, w763=obj.ii*0.5, w764=obj.jj*0.5, w765=obj.kk*0.5, w766=obj.ll*0.5, w767=obj.mm*0.5, w768=obj.nn*0.5, w769=obj.oo*0.5, w770=obj.pp*0.5, w771=obj.qq*0.5, w772=obj.rr*0.5, w773=obj.ss*0.5, w774=obj.tt*0.5, w775=obj.uu*0.5, w776=obj.vv*0.5, w777=obj.ww*0.5, w778=obj.xx*0.5, w779=obj.yy*0.5, w780=obj.zz*0.5, w781=obj.aa*0.5, w782=obj.bb*0.5, w783=obj.cc*0.5, w784=obj.dd*0.5, w785=obj.ee*0.5, w786=obj.ff*0.5, w787=obj.gg*0.5, w788=obj.hh*0.5, w789=obj.ii*0.5, w790=obj.jj*0.5, w791=obj.kk*0.5, w792=obj.ll*0.5, w793=obj.mm*0.5, w794=obj.nn*0.5, w795=obj.oo*0.5, w796=obj.pp*0.5, w797=obj.qq*0.5, w798=obj.rr*0.5, w799=obj.ss*0.5, w800=obj.tt*0.5, w801=obj.uu*0.5, w802=obj.vv*0.5, w803=obj.ww*0.5, w804=obj.xx*0.5, w805=obj.yy*0.5, w806=obj.zz*0.5, w807=obj.aa*0.5, w808=obj.bb*0.5, w809=obj.cc*0.5, w810=obj.dd*0.5, w811=obj.ee*0.5, w812=obj.ff*0.5, w813=obj.gg*0.5, w814=obj.hh*0.5, w815=obj.ii*0.5, w816=obj.jj*0.5, w817=obj.kk*0.5, w818=obj.ll*0.5, w819=obj.mm*0.5, w820=obj.nn*0.5, w821=obj.oo*0.5, w822=obj.pp*0.5, w823=obj.qq*0.5, w824=obj.rr*0.5, w825=obj.ss*0.5, w826=obj.tt*0.5, w827=obj.uu*0.5, w828=obj.vv*0.5, w829=obj.ww*0.5, w830=obj.xx*0.5, w831=obj.yy*0.5, w832=obj.zz*0.5, w833=obj.aa*0.5, w834=obj.bb*0.5, w835=obj.cc*0.5, w836=obj.dd*0.5, w837=obj.ee*0.5, w838=obj.ff*0.5, w839=obj.gg*0.5, w840=obj.hh*0.5, w841=obj.ii*0.5, w842=obj.jj*0.5, w843=obj.kk*0.5, w844=obj.ll*0.5, w845=obj.mm*0.5, w846=obj.nn*0.5, w847=obj.oo*0.5, w848=obj.pp*0.5, w849=obj.qq*0.5, w850=obj.rr*0.5, w851=obj.ss*0.5, w852=obj.tt*0.5, w853=obj.uu*0.5, w854=obj.vv*0.5, w855=obj.ww*0.5, w856=obj.xx*0.5, w857=obj.yy*0.5, w858=obj.zz*0.5, w859=obj.aa*0.5, w860=obj.bb*0.5, w861=obj.cc*0.5, w862=obj.dd*0.5, w863=obj.ee*0.5, w864=obj.ff*0.5, w865=obj.gg*0.5, w866=obj.hh*0.5, w867=obj.ii*0.5, w868=obj.jj*0.5, w869=obj.kk*0.5, w870=obj.ll*0.5, w871=obj.mm*0.5, w872=obj.nn*0.5, w873=obj.oo*0.5, w874=obj.pp*0.5, w875=obj.qq*0.5, w876=obj.rr*0.5, w877=obj.ss*0.5, w878=obj.tt*0.5, w879=obj.uu*0.5, w880=obj.vv*0.5, w881=obj.ww*0.5, w882=obj.xx*0.5, w883=obj.yy*0.5, w884=obj.zz*0.5, w885=obj.aa*0.5, w886=obj.bb*0.5, w887=obj.cc*0.5, w888=obj.dd*0.5, w889=obj.ee*0.5, w890=obj.ff*0.5, w891=obj.gg*0.5, w892=obj.hh*0.5, w893=obj.ii*0.5, w894=obj.jj*0.5, w895=obj.kk*0.5, w896=obj.ll*0.5, w897=obj.mm*0.5, w898=obj.nn*0.5, w899=obj.oo*0.5, w900=obj.pp*0.5, w901=obj.qq*0.5, w902=obj.rr*0.5, w903=obj.ss*0.5, w904=obj.tt*0.5, w905=obj.uu*0.5, w906=obj.vv*0.5, w907=obj.ww*0.5, w908=obj.xx*0.5, w909=obj.yy*0.5, w910=obj.zz*0.5, w911=obj.aa*0.5, w912=obj.bb*0.5, w913=obj.cc*0.5, w914=obj.dd*0.5, w915=obj.ee*0.5, w916=obj.ff*0.5, w917=obj.gg*0.5, w918=obj.hh*0.5, w919=obj.ii*0.5, w920=obj.jj*0.5, w921=obj.kk*0.5, w922=obj.ll*0.5, w923=obj.mm*0.5, w924=obj.nn*0.5, w925=obj.oo*0.5, w926=obj.pp*0.5, w927=obj.qq*0.5, w928=obj.rr*0.5, w929=obj.ss*0.5, w930=obj.tt*0.5, w931=obj.uu*0.5, w932=obj.vv*0.5, w933=obj.ww*0.5, w934=obj.xx*0.5, w935=obj.yy*0.5, w936=obj.zz*0.5, w937=obj.aa*0.5, w938=obj.bb*0.5, w939=obj.cc*0.5, w940=obj.dd*0.5, w941=obj.ee*0.5, w942=obj.ff*0.5, w943=obj.gg*0.5, w944=obj.hh*0.5, w945=obj.ii*0.5, w946=obj.jj*0.5, w947=obj.kk*0.5, w948=obj.ll*0.5, w949=obj.mm*0.5, w950=obj.nn*0.5, w951=obj.oo*0.5, w952=obj.pp*0.5, w953=obj.qq*0.5, w954=obj.rr*0.5, w955=obj.ss*0.5, w956=obj.tt*0.5, w957=obj.uu*0.5, w958=obj.vv*0.5, w959=obj.ww*0.5, w960=obj.xx*0.5, w961=obj.yy*0.5, w962=obj.zz*0.5, w963=obj.aa*0.5, w964=obj.bb*0.5, w965=obj.cc*0.5, w966=obj.dd*0.5, w967=obj.ee*0.5, w968=obj.ff*0.5, w969=obj.gg*0.5, w970=obj.hh*0.5, w971=obj.ii*0.5, w972=obj.jj*0.5, w973=obj.kk*0.5, w974=obj.ll*0.5, w975=obj.mm*0.5, w976=obj.nn*0.5, w977=obj.oo*0.5, w978=obj.pp*0.5, w979=obj.qq*0.5, w980=obj.rr*0.5, w981=obj.ss*0.5, w982=obj.tt*0.5, w983=obj.uu*0.5, w984=obj.vv*0.5, w985=obj.ww*0.5, w986=obj.xx*0.5, w987=obj.yy*0.5, w988=obj.zz*0.5, w989=obj.aa*0.5, w990=obj.bb*0.5, w991=obj.cc*0.5, w992=obj.dd*0.5, w993=obj.ee*0.5, w994=obj.ff*0.5, w995=obj.gg*0.5, w996=obj.hh*0.5, w997=obj.ii*0.5, w998=obj.jj*0.5, w999=obj.kk*0.5, w1000=obj.ll*0.5, w1001=obj.mm*0.5, w1002=obj.nn*0.5, w1003=obj.oo*0.5, w1004=obj.pp*0.5, w1005=obj.qq*0.5, w1006=obj.rr*0.5, w1007=obj.ss*0.5, w1008=obj.tt*0.5, w1009=obj.uu*0.5, w1010=obj.vv*0.5, w1011=obj.ww*0.5, w1012=obj.xx*0.5, w1013=obj.yy*0.5, w1014=obj.zz*0.5, w1015=obj.aa*0.5, w1016=obj.bb*0.5, w1017=obj.cc*0.5, w1018=obj.dd*0.5, w1019=obj.ee*0.5, w1020=obj.ff*0.5, w1021=obj.gg*0.5, w1022=obj.hh*0.5, w1023=obj.ii*0.5, w1024=obj.jj*0.5, w1025=obj.kk*0.5, w1026=obj.ll*0.5, w1027=obj.mm*0.5, w1028=obj.nn*0.5, w1029=obj.oo*0.5, w1030=obj.pp*0.5, w1031=obj.qq*0.5, w1032=obj.rr*0.5, w1033=obj.ss*0.5, w1034=obj.tt*0.5, w1035=obj.uu*0.5, w1036=obj.vv*0.5, w1037=obj.ww*0.5, w1038=obj.xx*0.5, w1039=obj.yy*0.5, w1040=obj.zz*0.5, w1041=obj.aa*0.
```

What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
    double t; // distance to intersection
    int id=0; // id of intersected object
    if (!intersect(r, t, id)) return Vec(); // if miss, return black
    const Sphere &obj = spheres[id]; // the hit object
    Vec x=r.o+r.d*t, n=(x-obj.p).norm(), nl=n.dot(r.d)<0?n:n*-1, f=obj.c;
    double p = f.x>f.y && f.x>f.z ? f.x : f.y>f.z ? f.y : f.z; // max refl
    if (++depth>5) if (erand48(Xi)<p) f=f*(1/p); else return obj.e; //R.R.
    if (obj.refl == DIFF){ // Ideal DIFFUSE reflection
        double r1=2*M_PI*erand48(Xi), r2=erand48(Xi), r2s=sqrt(r2);
        Vec w=nl, u=((fabs(w.x)>.1?Vec(0,1):Vec(1))%w).norm(), v=w%u;
        Vec d = (u*cos(r1)*r2s + v*sin(r1)*r2s + w*sqrt(1-r2)).norm();
        return obj.e + f.mult(radiance(Ray(x,d),depth,Xi));
    } else if (obj.refl == SPEC) // Ideal SPECULAR reflection
        return obj.e + f.mult(radiance(Ray(x,r.d-n*2*n.dot(r.d)),depth,Xi));
    Ray reflRay(x, r.d-n*2*n.dot(r.d)); // Ideal dielectric REFRACTION
    bool into = n.dot(nl)>0; // Ray from outside going in?
    double nc=1, nt=1.5, nnt=into?nc/nt:nt/nc, ddn=r.d.dot(nl), cos2t;
    if ((cos2t=1-nnt*nnt*(1-ddn*ddn))<0) // Total internal reflection
        return obj.e + f.mult(radiance(reflRay,depth,Xi));
    Vec tdir = (r.d*nnt - n*((into?-1):1)*(ddn*nnt+sqrt(cos2t))).norm();
    double a=nt-nc, b=nt+nc, R0=a*a/(b*b), c = 1-(into?-ddn:tdir.dot(n));
    double Re=R0+(1-R0)*c*c*c*c*c, Tr=1-Re, P=.25+.5*Re, RP=Re/P, TP=Tr/(1-P);
    return obj.e + f.mult(depth>2 ? (erand48(Xi)<P ? // Russian roulette
        radiance(reflRay,depth,Xi)*RP:radiance(Ray(x,tdir),depth,Xi)*TP) :
        radiance(reflRay,depth,Xi)*Re+radiance(Ray(x,tdir),depth,Xi)*Tr);
}
```

Optimizing smallpt

2020-11-04

What is smallpt anyway?

display the whole thing

What is smallpt anyway?

```
Vec radiance(const Ray &r, int depth, unsigned short *Xi){
    double t; // distance to intersection
    int id=0; // id of intersected object
    if (!intersect(r, t, id)) return Vec(); // if miss, return black
    const Sphere &obj = spheres[id]; // the hit object
    Vec x=r.o+r.d*t, n=(x-obj.p).norm(), nl=n.dot(r.d)<0?n:n*-1, f=obj.c;
    double p = f.x>f.y && f.x>f.z ? f.x : f.y>f.z ? f.y : f.z; // max refl
    if (++depth>5) if (erand48(Xi)<p) f=f*(1/p); else return obj.e; //R.R.
    if (obj.refl == DIFF){ // Ideal DIFFUSE reflection
        double r1=2*M_PI*erand48(Xi), r2=erand48(Xi), r2s=sqrt(r2);
        Vec w=nl, u=((fabs(w.x)>.1?Vec(0,1):Vec(1))%w).norm(), v=w%u;
        Vec d = (u*cos(r1)*r2s + v*sin(r1)*r2s + w*sqrt(1-r2)).norm();
        return obj.e + f.mult(radiance(Ray(x,d),depth,Xi));
    } else if (obj.refl == SPEC) // Ideal SPECULAR reflection
        return obj.e + f.mult(radiance(Ray(x,r.d-n*2*n.dot(r.d)),depth,Xi));
    Ray reflRay(x, r.d-n*2*n.dot(r.d)); // Ideal dielectric REFRACTION
    bool into = n.dot(nl)>0; // Ray from outside going in?
    double nc=1, nt=1.5, nnt=into?nc/nt:nt/nc, ddn=r.d.dot(nl), cos2t;
    if ((cos2t=1-nnt*nnt*(1-ddn*ddn))<0) // Total internal reflection
        return obj.e + f.mult(radiance(reflRay,depth,Xi));
    Vec tdir = (r.d*nnt - n*((into?-1):1)*(ddn*nnt+sqrt(cos2t))).norm();
    double a=nt-nc, b=nt+nc, R0=a*a/(b*b), c = 1-(into?-ddn:tdir.dot(n));
    double Re=R0+(1-R0)*c*c*c*c*c, Tr=1-Re, P=.25+.5*Re, RP=Re/P, TP=Tr/(1-P);
    return obj.e + f.mult(depth>2 ? (erand48(Xi)<P ? // Russian roulette
        radiance(reflRay,depth,Xi)*RP:radiance(Ray(x,tdir),depth,Xi)*TP) :
        radiance(reflRay,depth,Xi)*Re+radiance(Ray(x,tdir),depth,Xi)*Tr);
}
```

Initial Haskell Code: radiance (1×)

```
...
radiance :: Ray -> CInt -> Ptr Cushort -> IO Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
  (Nothing,_) -> return zerov
  (Just t,Sphere _r p e c refl) -> do
```

```
  continue f = case refl of
    DIFF -> do
```

```
      radiance
```

```
  SPEC -> do
```

```
      rad <- radiance
```

```
  REFR -> do
```

```
  if
  then do
    rad <- radiance reflRay depth' xi
```

Optimizing smallpt

2020-11-04

Initial Haskell Code: radiance (1×)

show the same thing, this time in haskell

Initial Haskell Code: radiance (1×)

```
radiance :: Ray -> CInt -> Ptr Cushort -> IO Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
  (Nothing,_) -> return zerov
  (Just t,Sphere _r p e c refl) -> do
```

```
  continue f = case refl of
    DIFF -> do
```

```
      radiance
```

```
  SPEC -> do
```

```
      rad <- radiance
```

```
  REFR -> do
```

```
  if
  then do
    rad <- radiance reflRay depth' xi
```

Initial Haskell Code: radiance (1×)

```
...
radiance :: Ray -> CInt -> Ptr CUSHort -> IO Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
  (Nothing,_) -> return zeroV
  (Just t,Sphere _r p e c refl) -> do
    let x = o `addv` (d `mulvs` t)
        n = norm $ x `subv` p
        nl = if n `dot` d < 0 then n else n `mulvs` (-1)
        pr = maxv c
        depth' = depth + 1
        continue f = case refl of
          DIFF -> do
            r1 <- ((2*pi)* `fmap` erand48 xi
            r2 <- erand48 xi
            let r2s = sqrt r2
                w@(Vec wx _ _) = nl
                u = norm $ (if abs wx > 0.1 then (Vec 0 1 0) else (Vec 1 0 0)) `cross` w
                v = w `cross` u
                d' = norm $ (u`mulvs`(cos r1*r2s)) `addv` (v`mulvs`(sin r1*r2s)) `addv` (w`mulvs`sqrt (1-r2))
            rad <- radiance (Ray x d') depth' xi
            return $ e `addv` (f `mulv` rad)
          SPEC -> do
            let d' = d `subv` (n `mulvs` (2 * (n`dot`d)))
            rad <- radiance (Ray x d') depth' xi
            return $ e `addv` (f `mulv` rad)
          REFR -> do
            let reflRay = Ray x (d `subv` (n `mulvs` (2* n`dot`d))) -- Ideal dielectric REFRACTION
                into = n`dot`nl > 0 -- Ray from outside going in?
                nc = 1
                nt = 1.5
                nnt = if into then nc/nt else nt/nc
                ddn = d`dot`nl
                cos2t = 1-nnt*nnt*(1-ddn*ddn)
            if cos2t<0 -- Total internal reflection
            then do
              rad <- radiance reflRay depth' xi
            ...

    if depth'>5
    ...
```

Optimizing smallpt

2020-11-04

Initial Haskell Code: radiance (1×)

full code



Initial Haskell Code: Entry point (1×)

```
smallpt :: Int -> Int -> Int -> IO ()
smallpt w h nsamps = do
  ...
  c <- VM.replicate (w * h) 0
  allocArray 3 \xi ->
    flip mapM_ [0..h-1] $ \y -> do
      writeXi xi y
      for_ [0..w-1] \x -> do
        let i = (h-y-1) * w + x
        for_ [0..1] \sy -> do
          for_ [0..1] \sx -> do
            r <- newIORef 0
            for_ [0..samps-1] \_s -> do
              r1 <- (2*) <$> erand48 xi
              ...
              rad <- radiance (Ray (org+d.*140) (norm d)) 0 xi
              ...
              modifyIORef r (+ rad .* recip (fromIntegral samps))
            ci <- VM.unsafeRead c i
            Vec rr rg rb <- readIORef r
            VM.unsafeWrite c i $ ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25
```

Optimizing smallpt

2020-11-04

Initial Haskell Code: Entry point (1×)

```
Initial Haskell Code: Entry point (1×)

smallpt :: Int -> Int -> Int -> IO ()
smallpt w h nsamps = do
  ...
  c <- VM.replicate (w * h) 0
  allocArray 3 \xi ->
    flip mapM_ [0..h-1] $ \y -> do
      writeXi xi y
      for_ [0..w-1] \x -> do
        let i = (h-y-1) * w + x
        for_ [0..1] \sy -> do
          for_ [0..1] \sx -> do
            r <- newIORef 0
            for_ [0..samps-1] \_s -> do
              r1 <- (2*) <$> erand48 xi
              ...
              rad <- radiance (Ray (org+d.*140) (norm d)) 0 xi
              ...
              modifyIORef r (+ rad .* recip (fromIntegral samps))
            ci <- VM.unsafeRead c i
            Vec rr rg rb <- readIORef r
            VM.unsafeWrite c i $ ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25
```

Initial Haskell Code: File I/O (1×)

```
withFile "image.ppm" WriteMode $ \hdl -> do
  hPrintf hdl "P3\n%d %d\n%d\n" w h (255::Int)
  flip mapM_ [0..w*h-1] \i -> do
    Vec r g b <- VM.unsafeRead c i
    hPrintf hdl "%d %d %d " (toInt r) (toInt g) (toInt b)
```

Optimizing smallpt

2020-11-04

Initial Haskell Code: File I/O (1×)

```
withFile "image.ppm" WriteMode $ \hdl -> do
  hPrintf hdl "P3\n%d %d\n%d\n" w h (255::Int)
  flip mapM_ [0..w*h-1] \i -> do
    Vec r g b <- VM.unsafeRead c i
    hPrintf hdl "%d %d %d " (toInt r) (toInt g) (toInt b)
```

Initial Haskell Code: RNG (1×)

```
foreign import ccall unsafe "erand48"  
  erand48 :: Ptr Cushort -> IO Double
```

Optimizing smallpt

2020-11-04

Initial Haskell Code: RNG (1×)

Initial Haskell Code: RNG (1×)

```
foreign import ccall unsafe "erand48"  
  erand48 :: Ptr Cushort -> IO Double
```


Restrict export list to 'main' (1.13x)

```
-module Main where
+module Main (main) where
```

Optimizing smallpt

2020-11-04

- Restrict export list to 'main' (1.13x)

Restrict export list to 'main' (1.13x)

```
-module Main where
+module Main (main) where
```

Knowing what functions are used how enables many optimizations that could otherwise would be detrimental or unsound (eg: changing signatures based on demands)

Mark entries of Ray and Sphere as UNPACK and Strict ($1.07\times$)

```
-data Ray = Ray Vec Vec -- origin, direction
+data Ray = Ray {-# UNPACK #-} !Vec {-# UNPACK #-} !Vec -- origin, direction
```

```
data Refl = DIFF | SPEC | REFR -- material types, used in radiance
```

```
-- radius, position, emission, color, reflection
-data Sphere = Sphere Double Vec Vec Vec !Refl
+data Sphere = Sphere {-# UNPACK #-} !Double
+                        {-# UNPACK #-} !Vec
+                        {-# UNPACK #-} !Vec
+                        {-# UNPACK #-} !Vec !Refl
```

- ▶ Removes laziness.
- ▶ Laziness implemented as unknown function call, which is an indirect jump, which is bad for performance

Optimizing smallpt

2020-11-04

- Mark entries of Ray and Sphere as UNPACK and Strict (1.07x)

[illegible]

This optimization removes indirection and laziness.

Optimizing smallpt

- Use a pattern synonym to unpack Refl in Sphere (1.07x)

```

--(4) LANGUAGE Fortran90C;
data Refl = SUFF | SPIN | NREF; -- material types, used in radiance
retype type Refl = Refl; -- material types, used in radiance
redefine SUFF,SPIN,NREF = Refl;
redefine SUFF = Refl 0
redefine SPIN = Refl 1
redefine NREF = Refl 2
--(5) COMPLETE COPY, SPIN, SUFF, NREF;
-- radius, position, emission, color, reflection
data Sphere = Sphere {# SPIN #} {# SUFF #} {# NREF #}
--(6) SPIN #} {# SUFF #} {# NREF #}
--(7) SPIN #} {# SUFF #} {# NREF #}
--(8) SPIN #} {# SUFF #} {# NREF #}
--(9) SPIN #} {# SUFF #} {# NREF #}

```

► Can be done using unbound sums. Is much uglier yet!

While we weren't able to unpack `Refl` in the last step because it is a sum type, we can if we use a trick modeled after an C Enum. The use of `COMPLETE` is unsafe here because other values could be constructed. For this small example though we rely on the fact that we create values of `Refl` via the patterns.

Change from maximum on a list to max (1.08×)

```
-maxv (Vec a b c) = maximum [a,b,c]
```

```
+maxv (Vec a b c) = max a (max b c)
```

```
@@ -84,7 +85,6 @@ radiance ray@(Ray o d) depth xi = case intersects ray of
```

```
    let x = o `addv` (d `mulvs` t)
```

```
    n = norm $ x `subv` p
```

```
    nl = if n `dot` d < 0 then n else n `mulvs` (-1)
```

```
-      pr = maxv c
```

```
    depth' = depth + 1
```

```
    continue f = case refl of
```

```
      DIFF -> do
```

```
@@ -140,6 +140,7 @@ radiance ray@(Ray o d) depth xi = case intersects ray of
```

```
    if depth'>5
```

```
    then do
```

```
      er <- erand48 xi
```

```
+      let !pr = maxv c
```

Optimizing smallpt

2020-11-04

Change from maximum on a list to max (1.08×)

```
-maxv (Vec a b c) = maximum [a,b,c]
+maxv (Vec a b c) = max a (max b c)
@@ -84,7 +85,6 @@ radiance ray@(Ray o d) depth xi = case intersects ray of
    let x = o `addv` (d `mulvs` t)
    n = norm $ x `subv` p
    nl = if n `dot` d < 0 then n else n `mulvs` (-1)
-    pr = maxv c
    depth' = depth + 1
    continue f = case refl of
      DIFF -> do
        er <- erand48 xi
        let !pr = maxv c
```

finicky optimization. GHC does not evaluate at compile time, making optimizations like these necessary

Convert erand48 to pure Haskell (1.09×)

```
-radiance :: Ray -> CInt -> Ptr CUSHort -> IO Vec
+radiance :: Ray -> Int -> IORef Word64 -> IO Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
  (Nothing,_) -> return zeroV
  (Just t,Sphere _r p e c refl) -> do
@@ -153,9 +153,8 @@ smallpt w h nsamps = do
    cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
    cy = norm (cx `cross` dir) `mulvs` 0.5135
    c <- VM.replicate (w * h) zeroV
-   allocaArray 3 $ \xi ->
-     flip mapM_ [0..h-1] $ \y -> do
+   xi <- newIORef 0
+   flip mapM_ [0..h-1] $ \y -> do
      writeXi xi y
```

Optimizing smallpt

2020-11-04

Convert erand48 to pure Haskell (1.09×)

Convert erand48 to pure Haskell (1.09×)

```
radiance :: Ray -> CInt -> Ptr CUSHort -> IO Vec
radiance :: Ray -> Int -> IORef Word64 -> IO Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
  (Nothing,_) -> return zeroV
  (Just t,Sphere _r p e c refl) -> do
@@ -153,9 +153,8 @@ smallpt w h nsamps = do
    cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
    cy = norm (cx `cross` dir) `mulvs` 0.5135
    c <- VM.replicate (w * h) zeroV
-   allocaArray 3 $ \xi ->
-     flip mapM_ [0..h-1] $ \y -> do
+   xi <- newIORef 0
+   flip mapM_ [0..h-1] $ \y -> do
      writeXi xi y
```

Remove mutability: Erand48 Monad

```
...
-erand48 :: IORef Word64 -> IO Double
-erand48 !t = do
-  r <- readIORef t
+data ET a = ET !Word64 !a deriving Functor
+newtype Erand48 a = Erand48 { runErand48' :: Word64 -> ET a } deriving Functor
+
+instance Applicative Erand48 where
+...
+instance Monad Erand48 where
+...
+runWithErand48 :: Int -> Erand48 a -> a
+runWithErand48 !y act =
+  let yw = fromIntegral y
+      prod = yw * yw * yw
+      ET _ !r = runErand48' act (prod `unsafeShiftL` 32) in r
+...
+erand48 :: Erand48 Double
+erand48 = Erand48 \ !r ->
+  let x' = 0x5deece66d * r + 0xb
+      d_word = 0x3ff0000000000000 .|. ((x' .&. 0xffffffff) `unsafeShiftL` 4)
+      d = castWord64ToDouble d_word - 1.0
```

Optimizing smallpt

2020-11-04

Remove mutability: Erand48 Monad

Remove mutability: Erand48 Monad

```
...
+erand48 :: IORef Word64 -> IO Double
+erand48 !t = do
+  r <- readIORef t
+data ET a = ET !Word64 !a deriving Functor
+newtype Erand48 a = Erand48 { runErand48' :: Word64 -> ET a } deriving Functor
+
+instance Applicative Erand48 where
+...
+instance Monad Erand48 where
+...
+runWithErand48 :: Int -> Erand48 a -> a
+runWithErand48 !y act =
+  let yw = fromIntegral y
+      prod = yw * yw * yw
+      ET _ !r = runErand48' act (prod `unsafeShiftL` 32) in r
+...
+erand48 :: Erand48 Double
+erand48 = Erand48 \ !r ->
+  let x' = 0x5deece66d * r + 0xb
+      d_word = 0x3ff0000000000000 .|. ((x' .&. 0xffffffff) `unsafeShiftL` 4)
+      d = castWord64ToDouble d_word - 1.0
```

Remove mutability: Using Erand48

```
-radiance :: Ray -> Int -> IORef Word64 -> IO Vec
-radiance ray@(Ray o d) depth xi = case intersects ray of
+radiance :: Ray -> Int -> Erand48 Vec
+radiance ray@(Ray o d) depth = case intersects ray of
...
-           r1 <- (2*pi*) <$> erand48 xi
-           r2 <- erand48 xi
+           r1 <- (2*pi*) <$> erand48
+           r2 <- erand48
...
-           then (* rp) <$> radiance reflRay depth' xi
-           else (* tp) <$> radiance (Ray x tdir) depth' xi
+           then (* rp) <$> radiance reflRay depth'
+           else (* tp) <$> radiance (Ray x tdir) depth'
...

```

Optimizing smallpt

2020-11-04

Remove mutability: Using Erand48

Remove mutability: Using Erand48

```
radiance :: Ray -> Int -> IORef Word64 -> IO Vec
radiance reflRay o d depth xi = case intersects ray of
radiance :: Ray -> Int -> Erand48 Vec
radiance reflRay o d depth = case intersects ray of
...
-           r1 <- (2*pi*) <$> erand48 xi
-           r2 <- erand48 xi
+           r1 <- (2*pi*) <$> erand48
+           r2 <- erand48
...
-           then (* rp) <$> radiance reflRay depth' xi
-           else (* tp) <$> radiance (Ray x tdir) depth' xi
+           then (* rp) <$> radiance reflRay depth'
+           else (* tp) <$> radiance (Ray x tdir) depth'
...

```

Removing mutability: eliminate IORef

Optimizing smallpt

2020-11-04

Removing mutability: eliminate IORef

Removing mutability: eliminate IORef

Set everything in smallpt to be strict (1.17x)

```
smallpt :: Int -> Int -> Int -> IO ()
smallpt w h nsamps = do
-   let samps = nsamps `div` 4
-   org = Vec 50 52 295.6
-   dir = norm $ Vec 0 (-0.042612) (-1)
-   cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
-   cy = norm (cx `cross` dir) `mulvs` 0.5135
+   let !samps = nsamps `div` 4
+   !org = Vec 50 52 295.6
+   !dir = norm $ Vec 0 (-0.042612) (-1)
+   !cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
+   !cy = norm (cx `cross` dir) `mulvs` 0.5135
...
-   r1 <- (2*) `fmap` erand48 xi
-   let dx = if r1<1 then sqrt r1-1 else 1-sqrt(2-r1)
-   r2 <- (2*) `fmap` erand48 xi
-   let dy = if r2<1 then sqrt r2-1 else 1-sqrt(2-r2)
-   d = ...
-   rad <- radiance (Ray (org`addv`(d`mulvs`140)) (norm d)) 0 xi
+   !r1 <- (2*) `fmap` erand48 xi
+   let !dx = if r1<1 then sqrt r1-1 else 1-sqrt(2-r1)
+   !r2 <- (2*) `fmap` erand48 xi
+   let !dy = if r2<1 then sqrt r2-1 else 1-sqrt(2-r2)
+   !d = ...
...
+   pure $! r + rad .* recip (fromIntegral samps)
+   pure $! ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25
```

Optimizing smallpt

2020-11-04

Set everything in smallpt to be strict (1.17x)

```
Set everything in smallpt to be strict (1.17x)
smallpt :: Int -> Int -> Int -> IO ()
smallpt w h nsamps = do
-   let samps = nsamps `div` 4
-   org = Vec 50 52 295.6
-   dir = norm $ Vec 0 (-0.042612) (-1)
-   cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
-   cy = norm (cx `cross` dir) `mulvs` 0.5135
+   let !samps = nsamps `div` 4
+   !org = Vec 50 52 295.6
+   !dir = norm $ Vec 0 (-0.042612) (-1)
+   !cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
+   !cy = norm (cx `cross` dir) `mulvs` 0.5135
...
-   r1 <- (2*) `fmap` erand48 xi
-   let dx = if r1<1 then sqrt r1-1 else 1-sqrt(2-r1)
-   r2 <- (2*) `fmap` erand48 xi
-   let dy = if r2<1 then sqrt r2-1 else 1-sqrt(2-r2)
-   d = ...
-   rad <- radiance (Ray (org`addv`(d`mulvs`140)) (norm d)) 0 xi
+   !r1 <- (2*) `fmap` erand48 xi
+   !dx = if r1<1 then sqrt r1-1 else 1-sqrt(2-r1)
+   !r2 <- (2*) `fmap` erand48 xi
+   !dy = if r2<1 then sqrt r2-1 else 1-sqrt(2-r2)
+   !d = ...
...
+   pure $! r + rad .* recip (fromIntegral samps)
+   pure $! ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25
```

This does make a difference but a small one. The question is which ones matter.

Reduce to only useful strictnesses in smallpt(1.17x)

```
- let !samps = nsamps `div` 4
- !org = Vec 50 52 295.6
- !dir = norm $ Vec 0 (-0.042612) (-1)
- !cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
- !cy = norm (cx `cross` dir) .* 0.5135
+ let samps = nsamps `div` 4
+ org = Vec 50 52 295.6
+ dir = norm $ Vec 0 (-0.042612) (-1)
+ cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
+ cy = norm (cx `cross` dir) .* 0.5135
...
- !r1 <- (2*) <$> erand48
+ r1 <- (2*) <$> erand48
...
- !r2 <- (2*) <$> erand48
+ r2 <- (2*) <$> erand48
...
- !rad <- radiance (Ray (org+d.*140) (norm d)) 0
+ rad <- radiance (Ray (org+d.*140) (norm d)) 0
...
- pure $! r + rad .* recip (fromIntegral samps)
- pure $! ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25
+ pure (r + rad .* recip (fromIntegral samps))
+ pure (ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25)
```

Optimizing smallpt

2020-11-04

Reduce to only useful strictnesses in smallpt(1.17x)

```
Reduce to only useful strictnesses in smallpt(1.17x)
- let tsamps = nsamps `div` 4
- !org = Vec 50 52 295.6
- !dir = norm $ Vec 0 (-0.042612) (-1)
- !cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
- !cy = norm (cx `cross` dir) .* 0.5135
+ let samps = nsamps `div` 4
+ org = Vec 50 52 295.6
+ dir = norm $ Vec 0 (-0.042612) (-1)
+ cx = Vec (fromIntegral w * 0.5135 / fromIntegral h) 0 0
+ cy = norm (cx `cross` dir) .* 0.5135
...
- !r1 <- (2*) <$> erand48
- !r2 <- (2*) <$> erand48
+ r1 <- (2*) <$> erand48
+ r2 <- (2*) <$> erand48
...
- !rad <- radiance (Ray (org+d.*140) (norm d)) 0
- !rad <- radiance (Ray (org+d.*140) (norm d)) 0
+ rad <- radiance (Ray (org+d.*140) (norm d)) 0
...
- pure $! r + rad .* recip (fromIntegral samps)
- pure $! ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25
+ pure (r + rad .* recip (fromIntegral samps))
+ pure (ci + Vec (clamp rr) (clamp rg) (clamp rb) .* 0.25)
```

Some inspection shows us which bang patters are actually carrying the weight. Most are unnecessary.

Use strictness strategically in entire project

```
- if det<0 then Nothing else f (b-sdet) (b+sdet)
- where op = p - o
-       eps = 1e-4
-       b = dot op d
-       det = b*b - dot op op + r*r
-       sdet = sqrt det
-       f a s = if a>eps then Just a else if s>eps then Just s else Nothing
+ if det<0
+ then Nothing
+ else
+   let !eps = 1e-4
+       !sdet = sqrt det
+       !a = b-sdet
+       !s = b+sdet
+   in if a>eps then Just a else if s>eps then Just s else Nothing
-   ...
```

TODO: maybe more interesting to talk about places where having strictness was ****not**** useful

Optimizing smallpt

2020-11-04

Use strictness strategically in entire project

Use strictness strategically in entire project

```
- if det<0 then Nothing else f (b-sdet) (b+sdet)
- where op = p - o
-       eps = 1e-4
-       b = dot op d
-       det = b*b - dot op op + r*r
-       sdet = sqrt det
-       f a s = if a>eps then Just a else if s>eps then Just s else Nothing
+ if det<0
+ then Nothing
+ else
+   let !eps = 1e-4
+       !sdet = sqrt det
+       !a = b-sdet
+       !s = b+sdet
+   in if a>eps then Just a else if s>eps then Just s else Nothing
-   ...
TODO: maybe more interesting to talk about places where having strictness was **not** useful
```

Remove Maybe from intersect(s) (1.32x)

```
-intersect :: Ray -> Sphere -> Maybe Double
+intersect :: Ray -> Sphere -> Double
intersect (Ray o d) (Sphere r p _e _c _refl) =
-   if det<0 then Nothing else f (b-sdet) (b+sdet)
+   if det<0 then (1/0.0) else f (b-sdet) (b+sdet)
  where op = p `subv` o
        eps = 1e-4
        b = op `dot` d
        det = b*b - (op `dot` op) + r*r
        sdet = sqrt det

-   f a s = if a>eps then Just a else if s>eps then Just s else Nothing
+   f a s = if a>eps then a else if s>eps then s else (1/0.0)

-intersects :: Ray -> (Maybe Double, Sphere)
+intersects :: Ray -> (Double, Sphere)
intersects ray = (k, s)
-   where (k,s) = foldl' f (Nothing,undefined) spheres
-         f (k',sp) s' = case (k',intersect ray s') of
-             (Nothing,Just x) -> (Just x,s')
-             (Just y,Just x) | x < y -> (Just x,s')
-             _ -> (k',sp)
+   where (k,s) = foldl' f (1/0.0,undefined) spheres
+         f (k', sp) s' = let !x = intersect ray s' in if x < k' then (x, s') else (k', sp)
```

```
radiance :: Ray -> Int -> STRefU s Word64 -> ST s Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
-   (Nothing,_) -> return zerov
-   (Just t,Sphere _r p e c refl) -> do
+   (t,_) | t == (1/0.0) -> return zerov
+   (t,Sphere _r p e c refl) -> do
```

Optimizing smallpt

2020-11-04

Remove Maybe from intersect(s) (1.32x)

Our innermost functions are of critical importance. Here we remove a Maybe which significantly reduces the boxing (which could have been mitigated with a StrictMaybe) and the cases. Since a Ray that fails to intersect something can be said to intersect at infinity, Double already actually covers the structure at play. This also reduces allocation.

```
Remove Maybe from intersect(s) (1.32x)
-intersect :: Ray -> Sphere -> Maybe Double
+intersect :: Ray -> Sphere -> Double
intersect (Ray o d) (Sphere r p _e _c _refl) =
-   if det<0 then Nothing else f (b-sdet) (b+sdet)
+   if det<0 then (1/0.0) else f (b-sdet) (b+sdet)
  where op = p `subv` o
        eps = 1e-4
        b = op `dot` d
        det = b*b - (op `dot` op) + r*r
        sdet = sqrt det

-   f a s = if a>eps then Just a else if s>eps then Just s else Nothing
+   f a s = if a>eps then a else if s>eps then s else (1/0.0)

-intersects :: Ray -> (Maybe Double, Sphere)
+intersects :: Ray -> (Double, Sphere)
intersects ray = (k, s)
-   where (k,s) = foldl' f (Nothing,undefined) spheres
-         f (k',sp) s' = case (k',intersect ray s') of
-             (Nothing,Just x) -> (Just x,s')
-             (Just y,Just x) | x < y -> (Just x,s')
-             _ -> (k',sp)
+   where (k,s) = foldl' f (1/0.0,undefined) spheres
+         f (k', sp) s' = let !x = intersect ray s' in if x < k' then (x, s') else (k', sp)

radiance :: Ray -> Int -> STRefU s Word64 -> ST s Vec
radiance ray@(Ray o d) depth xi = case intersects ray of
-   (Nothing,_) -> return zerov
-   (Just t,Sphere _r p e c refl) -> do
+   (t,_) | t == (1/0.0) -> return zerov
+   (t,Sphere _r p e c refl) -> do
```

Hand unroll the fold in intersects (1.35×)

```
intersects :: Ray -> (Double, Sphere)
-intersects ray = (k, s)
-  where (k,s) = foldl' f (1/0.0,undefined) spheres
+intersects ray =
+  f (... (f (f (intersect ray sphLeft, sphLeft) sphRight) ...)
+  where
+    f (k', sp) s' = let !x = intersect ray s' in if x < k' then (x, s') else (k', sp)

-spheres :: [Sphere]
-spheres = let s = Sphere ; z = zerov ; (.*) = mulvs ; v = Vec in
-  [ s 1e5 (v (1e5+1) 40.8 81.6)    z (v 0.75 0.25 0.25) DIFF --Left
-    , s 1e5 (v (-1e5+99) 40.8 81.6)  z (v 0.25 0.25 0.75) DIFF --Right
-  ...

+sphLeft, sphRight, ... :: Sphere
+sphLeft  = Sphere 1e5  (Vec (1e5+1) 40.8 81.6)    zerov (Vec 0.75 0.25 0.25) DIFF
+sphRight = Sphere 1e5  (Vec (-1e5+99) 40.8 81.6)  zerov (Vec 0.25 0.25 0.75) DIFF
+...
```

- Can also be realized using REWRITE rules.
- GHC does not attempt to partially evaluate.

Optimizing smallpt

2020-11-04

Hand unroll the fold in intersects (1.35×)

```
Hand unroll the fold in intersects (1.35×)

intersects :: Ray -> (Double, Sphere)
intersects ray = (k, s)
  where (k,s) = foldl' f (1/0.0,undefined) spheres
intersects ray =
  f (... (f (f (intersect ray sphLeft, sphLeft) sphRight) ...)
  where
    f (k', sp) s' = let !x = intersect ray s' in if x < k' then (x, s') else (k', sp)

spheres :: [Sphere]
spheres = let s = Sphere ; z = zerov ; (.*) = mulvs ; v = Vec in
  [ s 1e5 (v (1e5+1) 40.8 81.6)    z (v 0.75 0.25 0.25) DIFF --Left
    , s 1e5 (v (-1e5+99) 40.8 81.6)  z (v 0.25 0.25 0.75) DIFF --Right
  ...

sphLeft, sphRight, ... :: Sphere
sphLeft  = Sphere 1e5  (Vec (1e5+1) 40.8 81.6)    zerov (Vec 0.75 0.25 0.25) DIFF
sphRight = Sphere 1e5  (Vec (-1e5+99) 40.8 81.6)  zerov (Vec 0.25 0.25 0.75) DIFF
...
```

► Can also be realized using REWRITE rules.
► GHC does not attempt to partially evaluate.

This removes the list and this a potential level of indirections and branches. Sadly GHC did not do this for us even though the list was static.

Custom datatype for intersects parameter passing

```
-intersects :: Ray -> (Double, Sphere)
+data T = T !Double !Sphere
+
+intersects :: Ray -> T
+intersects ray =
+  f ( ... f (intersect ray sphLeft, sphLeft) sphRight) ... sphLite
+  where
+    f (k', sp) s' =
+      let !x = intersect ray s' in if x < k' then (x, s') else (k', sp)
+    f !(T k' sp) !s' =
+      let !x = intersect ray s' in if x < k' then T x s' else T k' sp
```

```
radiance :: Ray -> Int -> Erand48 Vec
radiance ray@(Ray o d) depth = case intersects ray of
-  (!t,_) | t == 1/0.0 -> return 0
-  (!t,!Sphere _r p e c refl) -> do
+  (T t _) | t == 1/0.0 -> return 0
+  (T t (Sphere _r p e c refl)) -> do
    let !x = o + d .* t
        !n = norm $ x - p
        !nl = if dot n d < 0 then n else negate n
```

Optimizing smallpt

2020-11-04

Custom datatype for intersects parameter passing

```
Custom datatype for intersects parameter passing

-intersects :: Ray -> (Double, Sphere)
+data T = T !Double !Sphere
+
+intersects :: Ray -> T
+intersects ray =
+  f ( ... f (intersect ray sphLeft, sphLeft) sphRight) ... sphLite
+  where
+    f (k', sp) s' =
+      let !x = intersect ray s' in if x < k' then (x, s') else (k', sp)
+    f !(T k' sp) !s' =
+      let !x = intersect ray s' in if x < k' then T x s' else T k' sp

radiance :: Ray -> Int -> Erand48 Vec
radiance ray@(Ray o d) depth = case intersects ray of
-  (!t,_) | t == 1/0.0 -> return 0
-  (!t,!Sphere _r p e c refl) -> do
+  (T t _) | t == 1/0.0 -> return 0
+  (T t (Sphere _r p e c refl)) -> do
    let !x = o + d .* t
        !n = norm $ x - p
        !nl = if dot n d < 0 then n else negate n
```

Its clear that f will consume these, so making it strict can have some benefit. The problem is the right level here. One's first take might be to set the parameters of 'f' to be strict, and that can have benefits, but it leaves a lot on the table. Moving to and unboxed tuple performs the computation faster but introduces copying. Out 'T' datatype here has the right balance, it makes it clear the data is strict but doesn't copy.

Optimize file writing

```
build-depends:
```

```
base >= 4.12 && < 4.15
```

```
+   , bytestring ^>= 0.11
```

```
-toInt :: Double -> Int
```

```
-toInt x = floor $ clamp x ** recip 2.2 * 255 + 0.5
```

```
+toInt :: Double -> BB.Builder
```

```
+toInt x = BB.intDec (floor (clamp x ** recip 2.2 * 255 + 0.5)) <> BB.char8 ' '
```

```
...
```

```
withFile "image.ppm" WriteMode $ \hdl -> do
```

```
-   hPrintf hdl "P3\n%d %d\n%d\n" w h (255::Int)
```

```
-   for_ img \(Vec r g b) -> do
```

```
-       hPrintf hdl "%d %d %d " (toInt r) (toInt g) (toInt b)
```

```
+   BB.hPutBuilder hdl $
```

```
+       BB.string8 "P3\n" <>
```

```
+       BB.intDec w <> BB.char8 ' ' <> BB.intDec h <> BB.char8 '\n' <>
```

```
+       BB.intDec 255 <> BB.char8 '\n' <>
```

```
+       (mconcat $ fmap \(Vec r g b) -> toInt r <> toInt g <> toInt b) img)
```

Optimizing smallpt

2020-11-04

Optimize file writing

Optimize file writing

```
build-depends:
  base >= 4.12 && < 4.15
+   , bytestring ^>= 0.11

-toInt :: Double -> Int
-toInt x = floor $ clamp x ** recip 2.2 * 255 + 0.5
+toInt :: Double -> BB.Builder
+toInt x = BB.intDec (floor (clamp x ** recip 2.2 * 255 + 0.5)) <> BB.char8 ' '
...
withFile "image.ppm" WriteMode $ \hdl -> do
-   hPrintf hdl "P3\n%d %d\n%d\n" w h (255::Int)
-   for_ img \(Vec r g b) -> do
-       hPrintf hdl "%d %d %d " (toInt r) (toInt g) (toInt b)
+   BB.hPutBuilder hdl $
+       BB.string8 "P3\n" <>
+       BB.intDec w <> BB.char8 ' ' <> BB.intDec h <> BB.char8 '\n' <>
+       BB.intDec 255 <> BB.char8 '\n' <>
+       (mconcat $ fmap \(Vec r g b) -> toInt r <> toInt g <> toInt b) img)
```

Strings are slow and the code was even using a slow conversion path to strings.

Use LLVM backend (1.87x)

```
+packages: .
+
+package smallpt-opt
+  ghc-options: -fllvm
```

Optimizing smallpt

2020-11-04


└ Use LLVM backend (1.87x)

Use LLVM backend (1.87x)

```
+packages: -
+
+package smallpt-xyz
+  ghc-options: -fllvm
```

The LLVM backend can pick up assembly level optimizations that GHC has missed. Some types of code runs worse with LLVM, but since this code is number crunching heavy, this does well.


The view from the mountaintop



Optimizing smallpt

2020-11-04

The view from the mountaintop



Takeaways

- ▶ Haskell can be fast, given some sensitivity to performance.
- ▶ Having performance leads to a faster baseline (unpacking, bang-patterns, `max`, LLVM by default, exporting `main`, ...)
- ▶ Some others (unrolling `f`) is more subtle.
- ▶ Accumulate optimizations to accrue performance wins.

Takeaways